



# How B.com avoids and deals with replication lag

Jean-François Gagné (System Engineer)

April 9, 2017 – MariaDB Developer Meeting 2017

# Booking.com

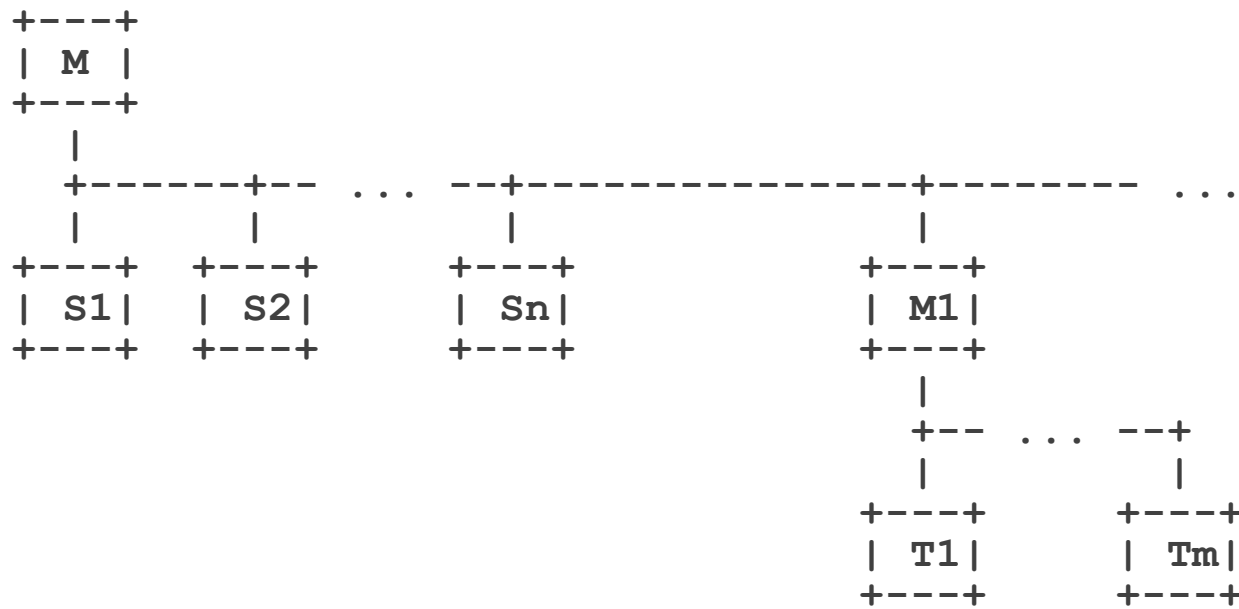
- Based in Amsterdam since 1996
- Online Hotel/Accommodation/Travel Agent (OTA):
  - +1.202.000 properties in 227 countries
  - +1.200.000 room nights reserved daily
  - +40 languages (website and customer service)
  - +13.000 people working in 187 offices worldwide
- Part of the Priceline Group
- And we use MySQL:
  - Thousands (1000s) of servers, ~90% replicating
  - >150 masters: ~30 >50 slaves & ~10 >100 slaves

# Session Summary

1. MySQL replication at Booking.com
2. Replication lag: what/how/why
3. Bad solutions to cope with lag
4. Booking.com solution to cope with lag
5. Improving Booking.com solution
6. Galera and Group Replication

# MySQL replication at Booking.com

- Typical Booking.com MySQL replication deployment:



# Why does lag happen ?

- In which condition can lag be experienced ?
  - Too many transactions for replication to keep up:  
capacity problem, fix by scaling (sharding, parallel replication, ...)
  - Long transactions:  
self induced, to fix by a developer in the application
  - Too aggressive “batch” workload on the master:  
optimize the batches or slow them down

# Lag consequences

- What are the consequences of lag ?
  - Stale reads on slaves (but this is not necessarily a problem)
- When do stale reads become a problem ?
  - A user changes his email address but still sees the old one
  - A hotel changes its inventory but still sees old availability
  - A user books a hotel but does not see it in his reservations

# Bad solution to cope with lag

- Bad solution #1: falling back to reading from the master
  - If slaves are lagging, maybe we should read from the master
  - This looks like an attractive solution to avoid stale reads
  - But this does not scale (why are you reading from slaves...)
  - This will cause a sudden load on the master (in case of lag)
  - And it might cause an outage on the master (and this would be bad)
- It might be better to fail a read than to fallback to (and kill) the master

# Bad solution to cope with lag (bis)

- Bad solution #2: retry on another slave
  - When reading from a slave: if lag, then retry on another slave
  - This scales better and is OK-ish (when few slaves are lagging)
  - But what happens if all slaves are lagging ?
  - Increased load (retries) can slowdown replication
  - This might overload the slaves and cause a good slave to start lagging
  - In the worst case, this might kill slaves and cause a domino effect
- Again: probably better to fail a read than to cause a bigger problem



# Coping with lag @ Booking.com

- Booking.com solution: “waypoint”
  - Creating a waypoint is similar to creating a “read view”
  - Waiting for a waypoint is similar to waiting for a slave to catch-up
- Booking.com implementation:
  - Table: db\_waypoint (a waypoint is a row in that table)
  - API function: commit\_wait(timeout) → (err\_code, waypoint)
    - INSERTs a waypoint and waits – until timeout – for its arrival on a slave
    - This is the same as creating a “read view” and “forcing” it on a slave
  - API function: waypoint\_wait(waypoint, timeout) → err\_code
    - Waits for a waypoint – until timeout – on a slave
    - This is the same as “waiting for a slave to catch-up”
  - Garbage collection: cleanup job that DELETES old waypoints

# Coping with lag @ Booking.com'

- Booking.com deployment:
  - Throttling batches:
    - use `commit_wait` with a high timeout
    - use “small” transactions (chunks of 100 to 1000 rows)
    - and sleep between chunks
  - Protect from stale reads after writing:
    - `commit_wait` with zero timeout
    - store the waypoint in web session
    - and `waypoint_wait` when reading

# Improving B.com waypoints

- The waypoint design and implementation still suits us.
- Sometime, we have a “fast” slave problem:
  - Throttling batches on a fast slave is sub-optimal
  - But this does not happen often in practice though
  - And it would be easy to fix: “find the slowest slave (or a slow slave)”
- But starting from scratch, we might do things differently:
  - Inserting, deleting and purging waypoint could be simplified
  - And we could get rid of the waypoint table

# Improving B.com waypoints'

- GTIDs as waypoint
  - Get the GTID of the last transaction:
    - last\_gtid session variable in MariaDB Server

From [https://mariadb.com/kb/en/mariadb/master\\_gtid\\_wait/](https://mariadb.com/kb/en/mariadb/master_gtid_wait/):

MASTER\_GTID\_WAIT() can also be used in client applications together with the last\_gtid session variable. This is useful in a read-scaleout replication setup, where the application writes to a single master but divides the reads out to a number of slaves to distribute the load. In such a setup, there is a risk that an application could first do an update on the master, and then a bit later do a read on a slave, and if the slave is not fast enough, the data read from the slave might not include the update just made, possibly confusing the application and/or the end-user. One way to avoid this is to request the value of last\_gtid on the master just after the update. Then before doing the read on the slave, do a MASTER\_GTID\_WAIT() on the value obtained from the master; this will ensure that the read is not performed until the slave has replicated sufficiently far for the update to have become visible.

# Improving B.com waypoints'

- GTIDs as waypoint:
  - Get the GTID of the last transaction :
    - last\_gtid session variable in MariaDB Server
    - gtid\_executed global variable in Oracle MySQL (get all executed GTIDs)
    - the last GTID can also be requested in the OK packet (only Oracle MySQL) (session\_track\_gtid variable and mysql\_session\_track\_get\_{first,next} API functions)
  - Waiting for GTID:
    - MASTER\_GTID\_WAIT in MariaDB Server
    - WAIT\_FOR\_EXECUTED\_GTID\_SET in Oracle MySQL
- But not portable (replicating from MySQL to MariaDB or vice-versa)

# Improving B.com waypoints''

- Binary log file and position as waypoint:
  - MASTER\_POS\_WAIT
  - However this breaks using intermediate masters
  - But it is OK with Binlog Servers<sup>[1]</sup>  
(in a Binlog Server deployment, the binlog file and position is a GTID)
  - But currently no way of getting file and position after committing

[1]: [https://blog.booking.com/abstracting\\_binlog\\_servers\\_and\\_mysql\\_master\\_promotion\\_wo\\_reconfiguring\\_slaves.html](https://blog.booking.com/abstracting_binlog_servers_and_mysql_master_promotion_wo_reconfiguring_slaves.html)

# Improving B.com waypoints'''

- Feature requests:
  - [Bug#84747](#): Expose last transaction GTID in a session variable.
  - [Bug#84748](#): Request transaction GTID in OK packet on COMMIT (without needing a round-trip).
  - [MDEV-11956](#): Get last\_gtid in OK packet.
  - [Bug#84779](#): Expose binlog file and position of last transaction.
  - [MDEV-11970](#): Expose binlog file and position of last transaction.

# Improving B.com waypoints''''

- Better solution for throttling:
  - Connecting to a (the right) slave is a hurdle
  - Having the information about slave state on the master would be great
  - A plugin exists for something close to that: semi-sync
  - Using this to track transaction execution on slaves would be great
  - This is the No-Slave-Left-Behind MariaDB Server Patch



# No-Slave-Left-Behind

- No-Slave-Left-Behind MariaDB Server patch<sup>[1]</sup>:
  - the semi-sync reply also reports SQL-thread position
  - transactions are kept in the master plugin until executed by one slave
  - the slave lag can be estimated from above
  - client-threads wait before commit until lag is acceptable
  - (Thanks Jonas Orelund and Google)
- This could easily be modified to implement `commit_wait`

[1]: <https://jira.mariadb.org/browse/MDEV-8112>

# Galera and Group Replication

- Group Replication allows to “read from the primary”
- Galera has a feature to protect against stale reads:
  - `wsrep_sync_wait=1` (or other values for other type of consistency)  
*the node blocks new queries while the database server catches up with all updates made in the cluster to the point where the check was begun*
- `wsrep_sync_wait=1` is very similar to checkpointing
- No-Slave-Left-Behind could also apply here:
  - No-Stale-Node-After-Write

# Links

- Booking.com:
  - <https://blog.booking.com/>
- MariaDB Server last\_gtid (thanks Kristian Nielsen for implementing this):
  - [https://mariadb.com/kb/en/mariadb/master\\_gtid\\_wait/](https://mariadb.com/kb/en/mariadb/master_gtid_wait/)
- Binlog Server:
  - [https://blog.booking.com/abstracting\\_binlog\\_servers\\_and\\_mysql\\_master\\_promotion\\_wo\\_reconfiguring\\_slaves.html](https://blog.booking.com/abstracting_binlog_servers_and_mysql_master_promotion_wo_reconfiguring_slaves.html)
- No-Slave-Left-Behind MariaDB Server patch:
  - <https://jira.mariadb.org/browse/MDEV-8112> (thanks Jonas Orelund and Google)

# Links'

- Pull request to extent Perl-DBI for reading GTID in OK packet:
  - <https://github.com/perl5-dbi/DBD-mysql/pull/77> (thanks Daniël van Eeden)
- Bug reports/Feature requests:
  - [Bug#84747: Expose last transaction GTID in a session variable.](#)
  - [Bug#84748: Request transaction GTID in OK packet on COMMIT \(without needing a round-trip\).](#)
  - [Bug#84779: Expose binlog file and position of last transaction.](#)
  - [MDEV-11956: Get last \\_gtid in OK packet.](#)
  - [MDEV-11970: Expose binlog file and position of last transaction.](#)

# Thanks

Jean-François Gagné  
jeanfrancois DOT gagne AT booking.com