



# How to Analyze and Tune MySQL Queries for Better Performance

Øystein Grøvlen  
Senior Principal Software Engineer  
MySQL Optimizer Team, Oracle  
March/April, 2016

**Please Stand By. This session will begin promptly at the time indicated on the agenda. Thank You.**

# Safe Harbor Statement

The following is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, and timing of any features or functionality described for Oracle's products remains at the sole discretion of Oracle.

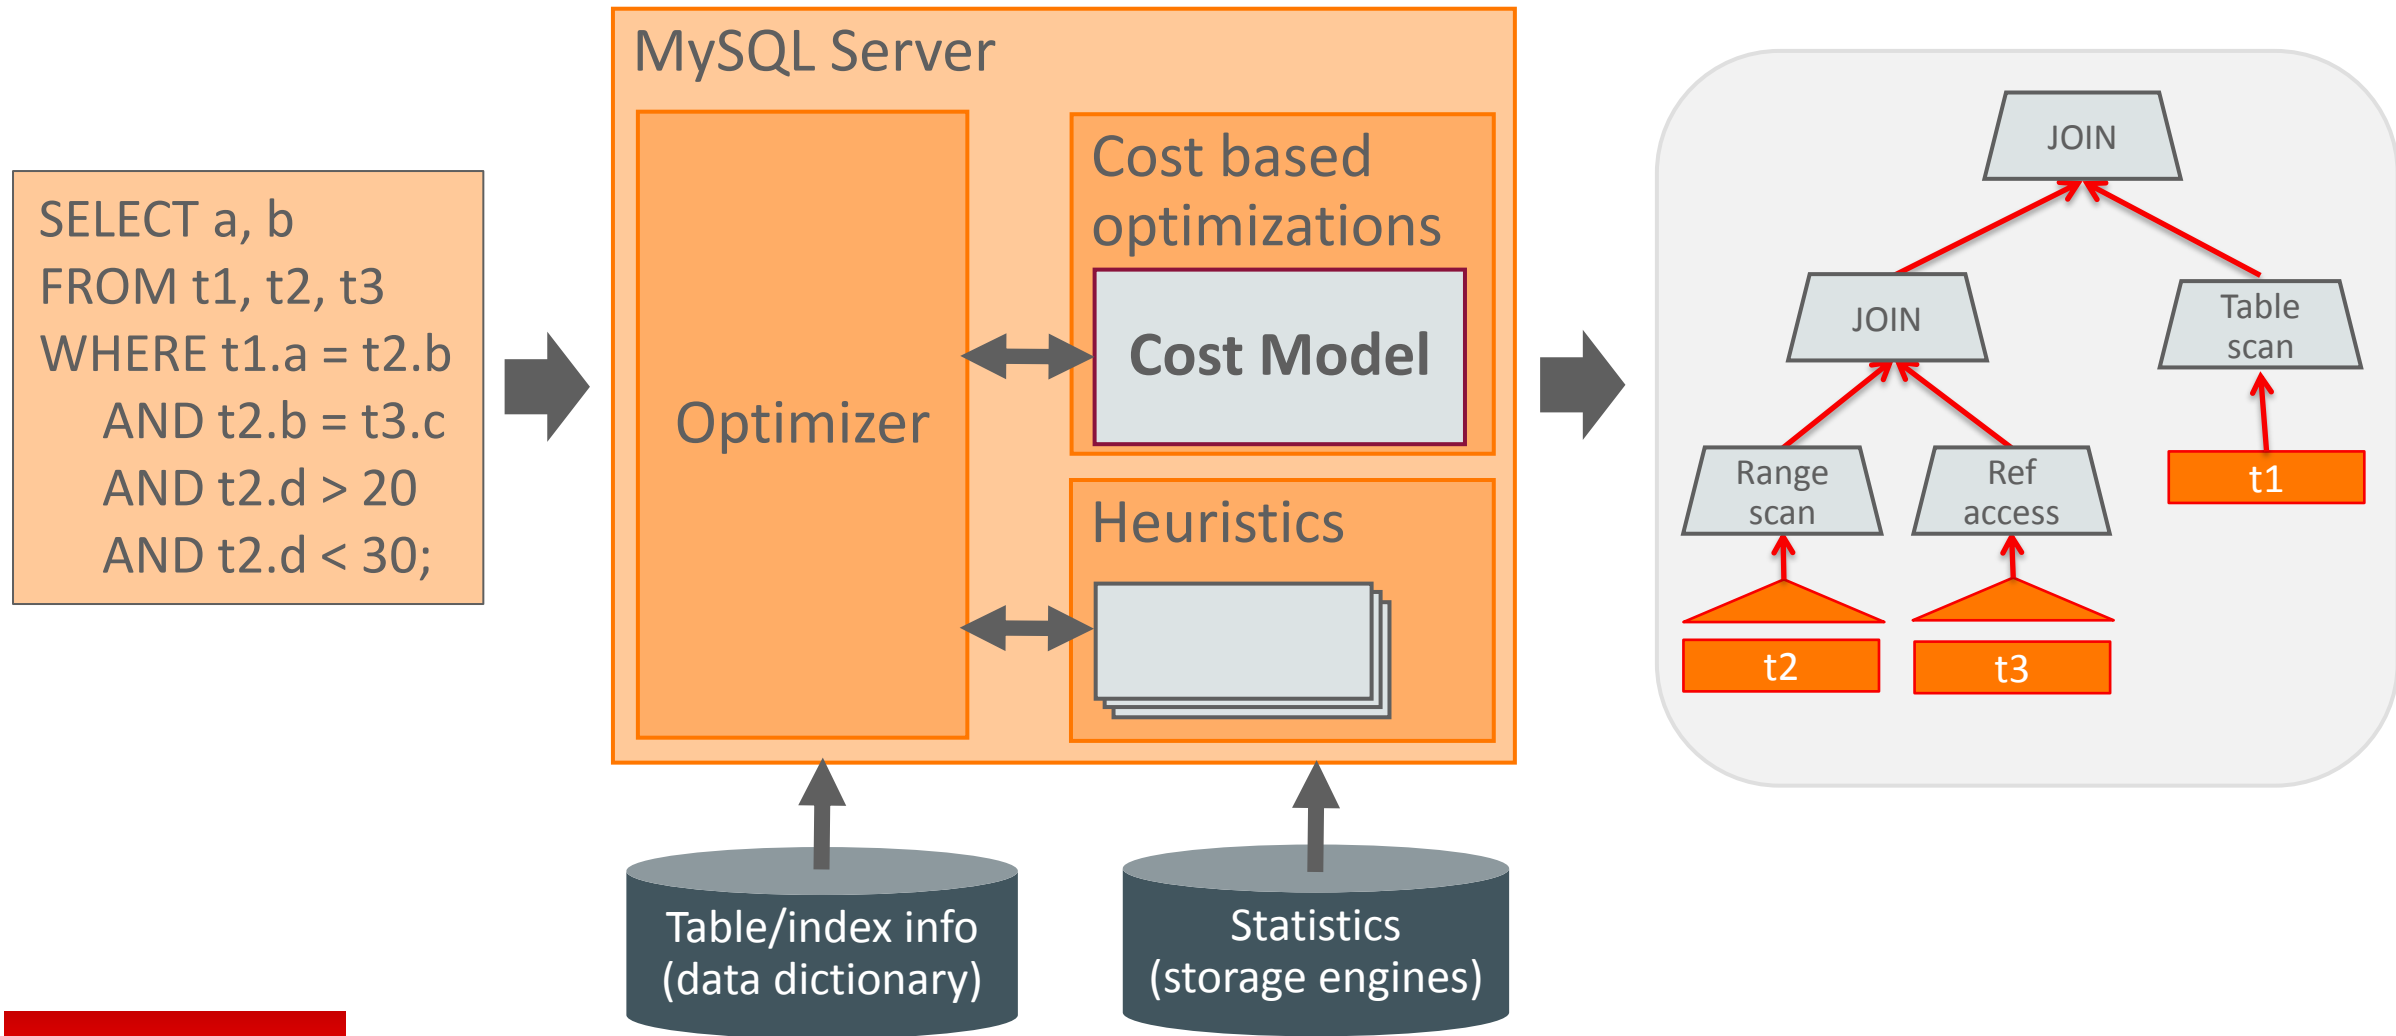
# Program Agenda

- 1 ➤ Cost-based query optimization in MySQL
- 2 ➤ Tools for monitoring, analyzing, and tuning queries
- 3 ➤ Data access and index selection
- 4 ➤ Join optimizer
- 5 ➤ Sorting
- 6 ➤ Influencing the optimizer

# Program Agenda

- 1 ➤ Cost-based query optimization in MySQL
- 2 ➤ Tools for monitoring, analyzing, and tuning queries
- 3 ➤ Data access and index selection
- 4 ➤ Join optimizer
- 5 ➤ Sorting
- 6 ➤ Influencing the optimizer

# MySQL Optimizer



# Cost-based Query Optimization

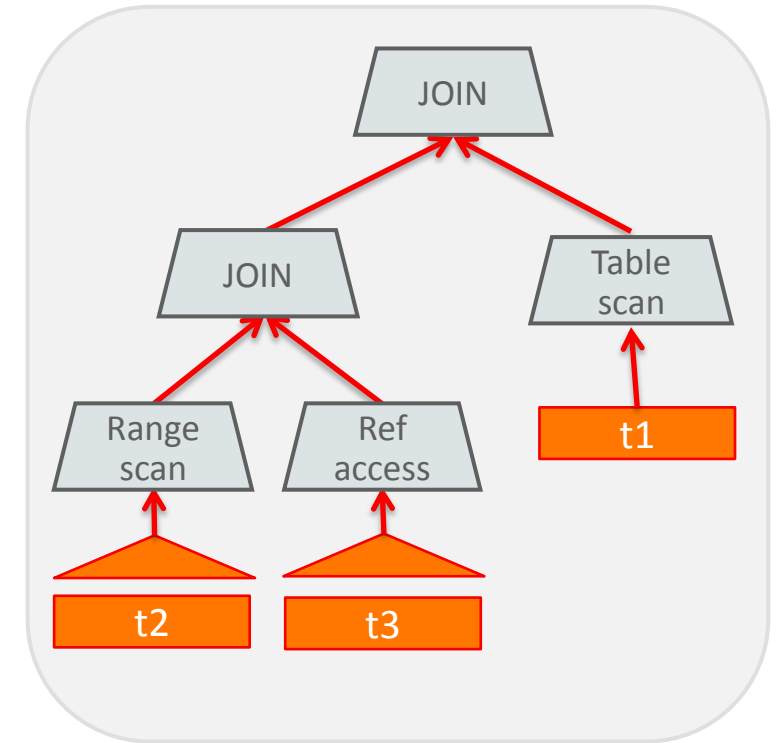
## General idea

- Assign cost to operations
- Assign cost to partial or alternative plans
- Search for plan with lowest cost
- Cost-based optimizations:

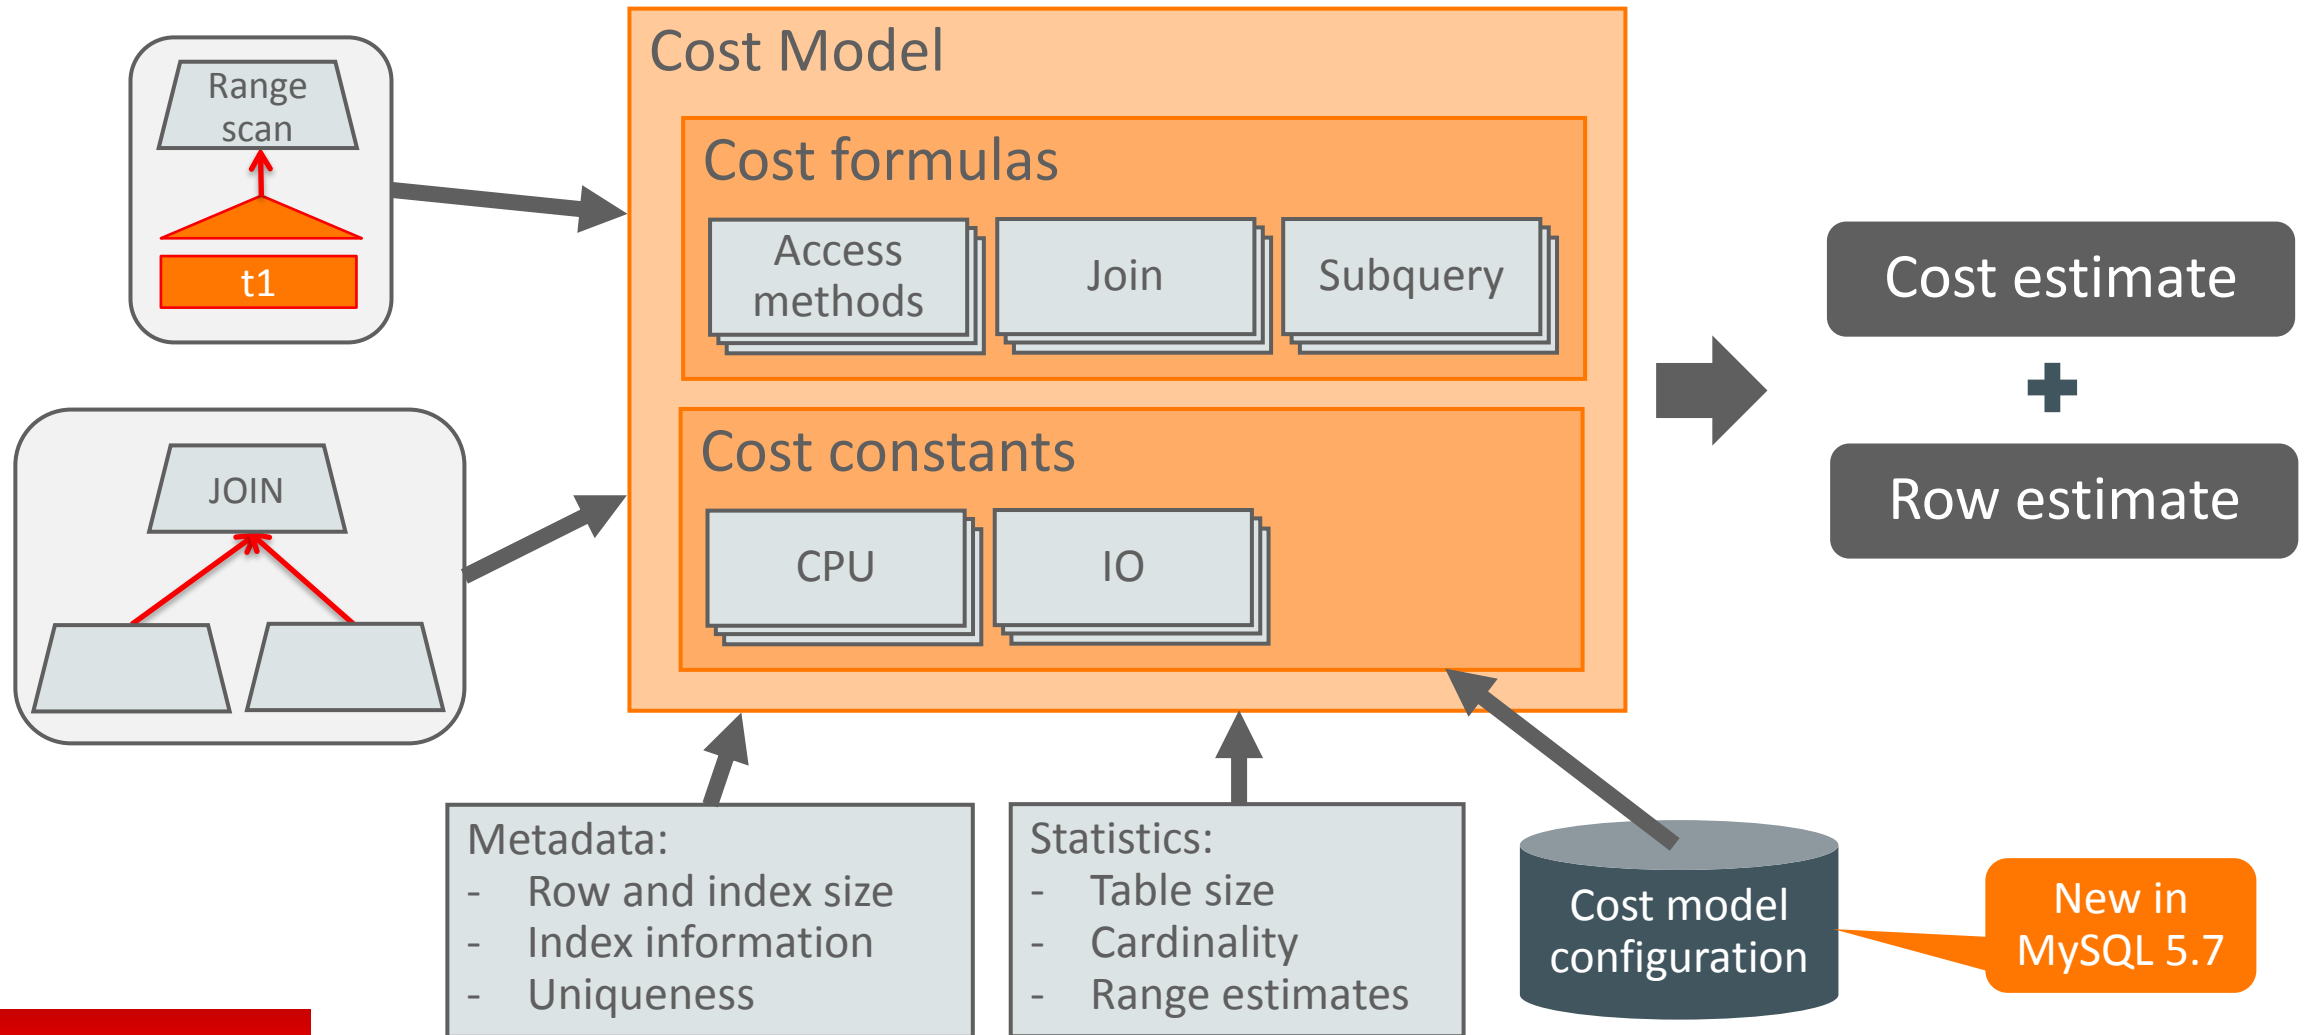
Access method

Join order

Subquery strategy



# Optimizer Cost Model




# Cost Model Example

```
SELECT SUM(o_totalprice) FROM orders
WHERE o_orderdate BETWEEN '1994-01-01' AND '1994-12-31';
```


## Table scan:

- IO-cost:  $\#pages \text{ in table} * IO\_BLOCK\_READ\_COST$
- CPU cost:  $\#rows * ROW\_EVALUATE\_COST$



## Range scan (on secondary index):

- IO-cost:  $\#rows\_in\_range * IO\_BLOCK\_READ\_COST$
- CPU cost:  $\#rows\_in\_range * ROW\_EVALUATE\_COST$





# Cost Model Example

EXPLAIN SELECT SUM(o\_totalprice) FROM orders  
WHERE o\_orderdate BETWEEN '1994-01-01' AND '1994-12-31';

id	select type	table	type	possible keys	key	key len	ref	rows	extra
1	SIMPLE	orders	ALL	i_o_orderdate	NULL	NULL	NULL	15000000	Using where

EXPLAIN SELECT SUM(o\_totalprice) FROM orders  
WHERE o\_orderdate BETWEEN '1994-01-01' AND '1994-06-30';

Id	select type	table	type	possible keys	key	key len	ref	rows	extra
1	SIMPLE	orders	range	i_o_orderdate	i_o_orderdate	4	NULL	2235118	Using index condition

# Cost Model Example: Optimizer Trace

## join\_optimization / row\_estimation / table : orders / range\_analysis

```
"table_scan": {
  "rows": 15000000,
  "cost": 3.12e6
} /* table_scan */,
"potential_range_indices": [
  {
    "index": "PRIMARY",
    "usable": false,
    "cause": "not_applicable"
  },
  {
    "index": "i_o_orderdate",
    "usable": true,
    "key_parts": [ "o_orderDATE", "o_orderkey" ]
  }
] /* potential_range_indices */,
...
```

```
"analyzing_range_alternatives": {
  "range_scan_alternatives": [
    {
      "index": "i_o_orderdate",
      "ranges": [ "1994-01-01 <= o_orderDATE <= 1994-12-31"
    ],
      "index_dives_for_eq_ranges": true,
      "rowid_ordered": false,
      "using_mrr": false,
      "index_only": false,
      "rows": 4489990,
      "cost": 5.39e6,
      "chosen": false,
      "cause": "cost"
    }
  ] /* range_scan_alternatives */,
  ...
} /* analyzing_range_alternatives */
```

# Cost Model vs Real World

## Measured Execution Times

	Data in Memory	Data on Disk	Data on SSD
Table scan	6.8 seconds	36 seconds	15 seconds
Index scan	<b>5.2 seconds</b>	2.5 hours	30 minutes

### Force Index Scan:

**SELECT SUM(o\_totalprice)**

**FROM orders **FORCE INDEX (i\_o\_orderdate)****

**WHERE o\_orderdate BETWEEN '1994-01-01' AND '1994-12-31';**

# Performance Schema

## Disk I/O

```
SELECT event_name, count_read, avg_timer_read/1000000000.0 "Avg Read Time (ms)",
       sum_number_of_bytes_read "Bytes Read"
FROM performance_schema.file_summary_by_event_name
WHERE event_name='wait/io/file/innodb/innodb_data_file';
```

### Table Scan

event_name	count_read	Avg Read Time (ms)	Bytes Read
wait/io/file/innodb/innodb_data_file	115769	0.0342	1896759296

### Index Scan

event_name	count_read	Avg Read Time (ms)	Bytes Read
wait/io/file/innodb/innodb_data_file	2188853	4.2094	35862167552

# Program Agenda

- 1 ➤ Cost-based query optimization in MySQL
- 2 ➤ Tools for monitoring, analyzing, and tuning queries
- 3 ➤ Data access and index selection
- 4 ➤ Join optimizer
- 5 ➤ Sorting
- 6 ➤ Influencing the optimizer

# Useful tools

- MySQL Enterprise Monitor (MEM), Query Analyzer
  - Commercial product
- Performance schema, MySQL sys schema
- EXPLAIN
  - Tabular EXPLAIN
  - Structured EXPLAIN (FORMAT=JSON)
  - Visual EXPLAIN (MySQL Workbench)
- Optimizer trace
- Slow log
- Status variables (SHOW STATUS LIKE 'Sort%')

# MySQL Enterprise Monitor, Query Analyzer



ORACLE MySQL Enterprise Monitor

1 2 0 0 1 admin Refresh: Every Minute

Browse Queries													
Show 25 entries Export data options... Showing 1 to 25 of 1,197 entries First Previous 1 2 3 4 5 Next Last													
Query	Database		Counts			QRTi	Latency (hh:mm:ss.ms)					Rows	
			Exec	Err	Warn		Total	Max	Avg	Locks	Avg History	Total	Examined
COMMIT (1)	mem		24,707	0	0	0.92	20:27.872	5.828	0.050	0.000		0	0
INSERT INTO `mem__quan`...o` ) , `hostTo` ) , ... (1)	mem		6,903	0	0	0.79	18:16.538	17.784	0.159	6.699		7,356	0
INSERT INTO `mem__quan`...on` = IF ( VALUES ( ... (1)	mem		6,985	0	0	0.86	10:54.856	8.940	0.094	8.301		7,332	0
INSERT INTO `mem__quan`...en` ) ) , `lastSeen` ) (1)	mem		7,025	37	0	1.00	3:11.791	11.220	0.027	4.436		13,947	0
INSERT INTO `mem__instr`...tency` ) , `latency` ) (1)	mem		740	0	0	0.77	1:48.459	8.745	0.147	0.147		1,386	0
SELECT `mysqlconne0` .....asProc18_1191_0_` , ... (1)	mem		974	0	0	0.90	1:01.829	12.359	0.063	0.353		974	974
SELECT `mysqlserve0` .....` . `hostCache` AS ... (1)	mem		1,801	0	0	0.93	53.661	13.351	0.030	0.636		1,801	1,801
INSERT INTO `mem__instr`...es` ) , `diskWrites` ) (1)	mem		26	0	0	0.79	44.838	14.954	1.725	0.070		26	0
UPDATE `mem__inventory`...n` = ? WHERE `hid` = ? (1)	mem		321	0	0	0.80	43.886	9.668	0.137	0.044		321	321
UPDATE `mem__config` . ... ? WHERE `user_id` = ? (1)	mem		321	0	0	0.71	40.788	10.607	0.127	0.038		321	321
CREATE TEMPORARY TABLE ... ( `id` INT8 NOT NULL ) (1)	mem		13	0	0	0.14	36.525	12.055	2.810	0.000		0	0
INSERT INTO `mem__instr`...nedTableDefinitions` ) (1)	mem		26	0	0	0.64	34.348	13.349	1.321	0.003		27	0
UPDATE `mem__inventory`...p` = ? WHERE `hid` = ? (1)	mem		416	0	0	0.81	33.469	13.060	0.080	0.042		416	416
INSERT INTO `mem__instr`...hed` ) , `notCached` ) (1)	mem		26	0	0	0.61	32.509	11.782	1.250	0.003		27	0
SELECT * FROM ( SELECT ...m_no_index_used` AS ... (1)	mem		14	0	0	0.00	29.832	9.602	2.131	0.008		7,459	227,050
INSERT INTO `mem__instr`...s` ) , `connections` ) (1)	mem		25	0	0	0.82	29.462	14.294	1.178	0.005		26	0
INSERT INTO `mem__instr`... ( `sent` ) , `sent` ) (1)	mem		25	0	0	0.82	28.991	14.332	1.160	0.003		26	0
SELECT `agent0` . `hid`.... `hid` = ? FOR UPDATE (1)	mem		909	0	0	0.96	28.632	6.964	0.031	0.158		909	909
UPDATE `mem__events` . ...me` = ? WHERE `id` = ? (1)	mem		395	0	0	0.92	27.544	10.061	0.070	0.038		395	395

Copyright © 2005, 2013, Oracle and/or its affiliates. All rights reserved.

3.0.1.7150 - Cerberus.local (192.168.1.67) - Sep 16, 2013 11:50:28 am (Up Since: 38 minutes ago) - About

# Query Analyzer Query Details

Canonical Query **Example Query** Explain Query Graphs

The query with the longest execution time during the Time Span (usually the slowest but not always).

## Sampled Query

[truncated](#) | [full](#) | [formatted](#)

```
SELECT
mysqlserve0.`hid` AS hid1124_0, mysqlserve0.`id` AS id2_1124_0,
mysqlserve0.`lastContact` AS lastCont3_1124_0,
mysqlserve0.`hasLastContact` AS hasLastC4_1124_0,
mysqlserve0.`startTime` AS startTime5_1124_0,
mysqlserve0.`hasStartTime` AS hasStart6_1124_0,
mysqlserve0.`timestamp` AS timestamp7_1124_0,
mysqlserve0.`capabilities` AS capabili8_1124_0,
mysqlserve0.`hasCapabilities` AS hasCapab9_1124_0,
mysqlserve0.`characterSet` AS charact10_1124_0,
mysqlserve0.`hasCharacterSet` AS hasChar11_1124_0,
mysqlserve0.`collation` AS collation12_1124_0,
mysqlserve0.`hasCollation` AS hasColl13_1124_0,
mysqlserve0.`connection` AS connection14_1124_0,
mysqlserve0.`hasConnection` AS hasConn15_1124_0,
mysqlserve0.`environment` AS environ16_1124_0,
```

## Execution Time

27,084 ms

## Date

Sep 16, 2013 1:07:17 PM

## User

service\_manager

## Thread ID

10,712

## From Host

localhost

## To Host

## Source Location

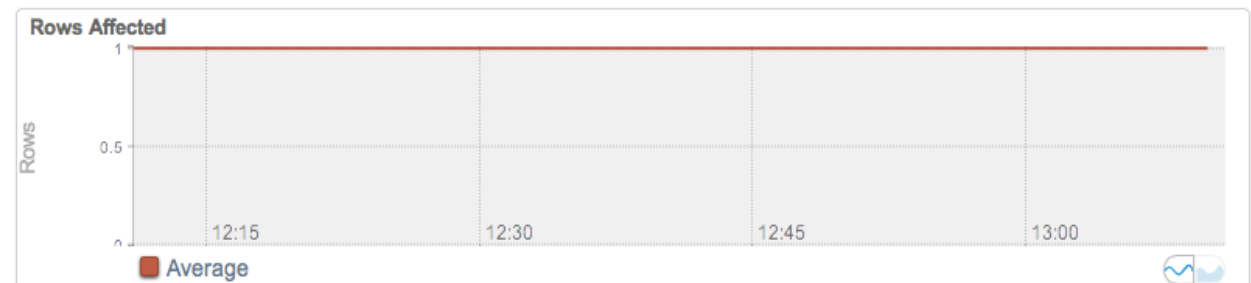


None found.

## Comments

None found.

Canonical Query Example Query Explain Query **Graphs**

Graphs of a query during the Time Span





# Performance Schema

## Some useful tables

- `events_statements_history`  
`events_statements_history_long`
  - Most recent statements executed
- `events_statements_summary_by_digest`
  - Summary for similar statements (same statement digest)
- `file_summary_by_event_name`
  - Interesting event: `wait/io/file/innodb/innodb_data_file`
- `table_io_waits_summary_by_table`  
`table_io_waits_summary_by_index_usage`
  - Statistics on storage engine access per table and index

# Performance Schema

## Statement events

- Tables:

`events_statements_current` (Current statement for each thread)

`events_statements_history` (10 most recent statements per thread)

`events_statements_history_long` (10000 most recent statements)

- Columns:

`THREAD_ID`, `EVENT_ID`, `END_EVENT_ID`, `EVENT_NAME`, `SOURCE`, `TIMER_START`, `TIMER_END`, `TIMER_WAIT`,  
`LOCK_TIME`, `SQL_TEXT`, `DIGEST`, `DIGEST_TEXT`, `CURRENT_SCHEMA`, `OBJECT_TYPE`, `OBJECT_SCHEMA`,  
`OBJECT_NAME`, `OBJECT_INSTANCE_BEGIN`, `MYSQL_ERRNO`, `RETURNED_SQLSTATE`, `MESSAGE_TEXT`, `ERRORS`,  
`WARNINGS`, `ROWS_AFFECTED`, `ROWS_SENT`, `ROWS_EXAMINED`, `CREATED_TMP_DISK_TABLES`,  
`CREATED_TMP_TABLES`, `SELECT_FULL_JOIN`, `SELECT_FULL_RANGE_JOIN`, `SELECT_RANGE`, `SELECT_RANGE_CHECK`,  
`SELECT_SCAN`, `SORT_MERGE_PASSES`, `SORT_RANGE`, `SORT_ROWS`, `SORT_SCAN`, `NO_INDEX_USED`,  
`NO_GOOD_INDEX_USED`, `NESTING_EVENT_ID`, `NESTING_EVENT_TYPE`

# Performance Schema

## Statement digest

- Normalization of queries to group statements that are similar to be grouped and summarized:

**SELECT \* FROM orders WHERE o\_custkey=10 AND o\_totalprice>20**

**SELECT \* FROM orders WHERE o\_custkey = 20 AND o\_totalprice > 100**

**➡ SELECT \* FROM orders WHERE o\_custkey = ? AND o\_totalprice > ?**

- **events\_statements\_summary\_by\_digest**

**DIGEST, DIGEST\_TEXT, COUNT\_STAR, SUM\_TIMER\_WAIT, MIN\_TIMER\_WAIT, AVG\_TIMER\_WAIT, MAX\_TIMER\_WAIT, SUM\_LOCK\_TIME, SUM\_ERRORS, SUM\_WARNINGS, SUM\_ROWS\_AFFECTED, SUM\_ROWS\_SENT, SUM\_ROWS\_EXAMINED, SUM\_CREATED\_TMP\_DISK\_TABLES, SUM\_CREATED\_TMP\_TABLES, SUM\_SELECT\_FULL\_JOIN, SUM\_SELECT\_FULL\_RANGE\_JOIN, SUM\_SELECT\_RANGE, SUM\_SELECT\_RANGE\_CHECK, SUM\_SELECT\_SCAN, SUM\_SORT\_MERGE\_PASSES, SUM\_SORT\_RANGE, SUM\_SORT\_ROWS, SUM\_SORT\_SCAN, SUM\_NO\_INDEX\_USED, SUM\_NO\_GOOD\_INDEX\_USED, FIRST\_SEEN, LAST\_SEEN**

# MySQL sys Schema

- A collection of views, procedures and functions, designed to make reading raw Performance Schema data easier
- Implements many common DBA and Developer use cases
  - File IO usage per user
  - Which indexes is never used?
  - Which queries use full table scans?
- Examples of very useful functions:
  - `format_time()` , `format_bytes()`, `format_statement()`
- **Included with MySQL 5.7**
- Bundled with MySQL Workbench

# MySQL sys Schema

## Example

**statement\_analysis:** Lists a normalized statement view with aggregated statistics, ordered by the total execution time per normalized statement

```
mysql> SELECT * FROM sys.statement_analysis LIMIT 1\G
```

```
***** 1. row *****
```

```
query: INSERT INTO `mem__quan`. `nor ... nDuration` = IF ( VALUES ( ...
```

```
db: mem
```

```
full_scan: 0
```

```
exec_count: 1110067
```

```
err_count: 0
```

```
warn_count: 0
```

```
total_latency: 1.93h
```

```
max_latency: 5.03 s
```

```
avg_latency: 6.27 ms
```

```
lock_latency: 00:18:29.18
```

```
rows_sent: 0
```

```
rows_sent_avg: 0
```

```
rows_examined: 0
```

```
rows_examined_avg: 0
```

```
tmp_tables: 0
```

```
tmp_disk_tables: 0
```

```
rows_sorted: 0
```

```
sort_merge_passes: 0
```

```
digest: d48316a218e95b1b8b72db5e6b177788!
```

```
first_seen: 2014-05-20 10:42:17
```

# EXPLAIN

## Understand the query plan

- Use **EXPLAIN** to print the final query plan:

```
EXPLAIN SELECT * FROM t1 JOIN t2 ON t1.a = t2.a WHERE b > 10 AND c > 10;
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	t1	NULL	range	idx1	idx1	4	NULL	12	33.33	Using where
1	SIMPLE	t2	NULL	ref	idx2	idx2	4	t1.a	1	100.00	NULL

- Explain for a running query (**New in MySQL 5.7**):

**EXPLAIN FOR CONNECTION** *connection\_id*;

# Structured EXPLAIN

- JSON format:

**EXPLAIN FORMAT=JSON SELECT ...**

- Contains more information:

- Used index parts
- Pushed index conditions
- Cost estimates
- Data estimates

Added in  
MySQL 5.7

```
EXPLAIN FORMAT=JSON
SELECT * FROM t1 WHERE b > 10 AND c > 10;
EXPLAIN
{
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "17.81"
    },
    "table": {
      "table_name": "t1",
      "access_type": "range",
      "possible_keys": [
        "idx1"
      ],
      "key": "idx1",
      "used_key_parts": [
        "b"
      ],
      "key_length": "4",
      "rows_examined_per_scan": 12,
      "rows_produced_per_join": 3,
      "filtered": "33.33",
      "index_condition": "(`test`.`t1`.`b` > 10)",
      "cost_info": {
        "read_cost": "17.01",
        "eval_cost": "0.80",
        "prefix_cost": "17.81",
        "data_read_per_join": "63"
      },
      "attached_condition": "(`test`.`t1`.`c` > 10)"
    }
  }
}
```

# Structured EXPLAIN

## Assigning Conditions to Tables

```
EXPLAIN FORMAT=JSON SELECT * FROM t1, t2
WHERE t1.a=t2.a AND t2.a=9 AND (NOT (t1.a > 10 OR t2.b >3) OR (t1.b=t2.b+7 AND t2.b = 5));
```

### EXPLAIN

```
{
  "query_block": {
    "select_id": 1,
    "nested_loop": [
      {
        "table": {
          "table_name": "t1",
          "access_type": "ALL",
          "rows": 10,
          "filtered": 100,
          "attached_condition": "(t1.a = 9)"
        } /* table */
      },

```

```
{
  "table": {
    "table_name": "t2",
    "access_type": "ALL",
    "rows": 10,
    "filtered": 100,
    "using_join_buffer": "Block Nested Loop",
    "attached_condition": "((t2.a = 9) and ((t2.b <= 3) or ((t2.b =
5) and (t1.b = 12))))"
  } /* table */
] /* nested_loop */
} /* query_block */
}
```



# Optimizer Trace: Query Plan Debugging

- EXPLAIN shows the selected plan
- Optimizer trace shows WHY the plan was selected

```
SET optimizer_trace= "enabled=on";
SELECT * FROM t1,t2 WHERE f1=1 AND f1=f2 AND f2>0;
SELECT trace FROM information_schema.optimizer_trace
INTO OUTFILE <filename> LINES TERMINATED BY ";
SET optimizer_trace="enabled=off";
```

QUERY	SELECT * FROM t1,t2 WHERE f1=1 AND f1=f2 AND f2>0;
TRACE	"steps": [ { "join_preparation": { "select#": 1,... } ... } ...]
MISSING_BYTES_BEYOND_MAX_MEM_SIZE	0
INSUFFICIENT_PRIVILEGES	0

# Program Agenda

- 1 ➤ Cost-based query optimization in MySQL
- 2 ➤ Tools for monitoring, analyzing, and tuning queries
- 3 ➤ Data access and index selection**
- 4 ➤ Join optimizer
- 5 ➤ Sorting
- 6 ➤ Influencing the optimizer

# Selecting Access Method

## Finding the optimal method to read data from storage engine

- For each table, find the best access method:
  - Check if the access method is useful
  - Estimate cost of using access method
  - Select the cheapest to be used
- Choice of access method is cost based

### Main access methods:

- Table scan
- Index scan
- Index look-up (ref access)
- Range scan
- Index merge
- Loose index scan

# Ref Access

## Single Table Queries

**EXPLAIN SELECT \* FROM customer WHERE c\_custkey = 570887;**

id	select type	table	type	possible keys	key	key len	ref	rows	extra
1	SIMPLE	customer	const	PRIMARY	PRIMARY	4	const	1	NULL

**EXPLAIN SELECT \* FROM orders WHERE o\_orderdate = '1992-09-12';**

id	select type	table	type	possible keys	key	key len	ref	rows	extra
1	SIMPLE	orders	ref	i_o_orderdate	i_o_orderdate	4	const	6271	NULL

# Ref Access

## Join Queries

**EXPLAIN SELECT \***

**FROM orders JOIN customer ON c\_custkey = o\_custkey**

**WHERE o\_orderdate = '1992-09-12';**

id	select type	table	type	possible keys	key	key len	ref	rows	extra
1	SIMPLE	orders	ref	i_o_orderdate, i_o_custkey	i_o_orderdate	4	const	6271	Using where
1	SIMPLE	customer	eq_ref	PRIMARY	PRIMARY	4	dbt3.orders. o_custkey	1	NULL

# Range Optimizer

- Goal: find the "minimal" ranges for each index that needs to be read


- Example:

**SELECT \* FROM t1 WHERE (key1 > 10 AND key1 < 20) AND key2 > 30**

- Range scan using INDEX(key1):



- Range scan using INDEX(key2):



# Range Optimizer: Case Study

## Why table scan?

```
SELECT * FROM orders
WHERE YEAR(o_orderdate) = 1997 AND MONTH(o_orderdate) = 5
      AND o_clerk = 'Clerk#000001866';
```

id	select type	table	type	possible keys	key	key len	ref	rows	extra
1	SIMPLE	orders	ALL	NULL	NULL	NULL	NULL	15000000	Using where

Index not considered

```
mysql> SELECT * FROM orders WHERE year(o_orderdate) = 1997 AND MONTH(...)
...
15 rows in set (8.91 sec)
```

# Range Optimizer: Case Study

Rewrite query to avoid functions on indexed columns

```
SELECT * FROM orders
WHERE o_orderdate BETWEEN '1997-05-01' AND '1997-05-31'
AND o_clerk = 'Clerk#000001866';
```

id	select type	table	type	possible keys	key	key len	ref	rows	extra
1	SIMPLE	orders	range	i_o_orderdate	i_o_orderdate	4	NULL	376352	Using index condition; Using where

```
mysql> SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' AND ...
...
15 rows in set (0.91 sec)
```



# Range Optimizer: Case Study

Adding another index

```
CREATE INDEX i_o_clerk ON orders(o_clerk);
```

```
SELECT * FROM orders
```

```
WHERE o_orderdate BETWEEN '1997-05-01' AND '1997-05-31'  
AND o_clerk = 'Clerk#000001866';
```


id	select type	table	type	possible keys	key	key len	ref	rows	extra
1	SIMPLE	orders	range	i_o_orderdate, i_o_clerk	i_o_clerk	16	NULL	1504	Using index condition; Using where

```
mysql> SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' AND ...  
...  
15 rows in set (0.01 sec)
```

# Range Access for Multi-Column Indexes


Example table with multi-part index

- Table:



- INDEX idx(a, b, c);

- Logical storage layout of index:




# Range Access for Multi-Column Indexes, cont

- Equality on 1<sup>st</sup> index column?
  - Can add condition on 2<sup>nd</sup> index column to range condition

- Example:

**SELECT \* from t1 WHERE a IN (10,11,13) AND (b=2 OR b=4)**




- Resulting range scan:



# Range Access for Multi-Column Indexes, cont

- Non-Equality on 1<sup>st</sup> index column:
  - Can **NOT** add condition on 2<sup>nd</sup> index column to range condition
- Example:

**SELECT \* from t1 WHERE a > 10 AND a < 13 AND (b=2 OR b=4)**



- Resulting range scan:



# Range Optimizer: Case Study

Create multi-column index

**CREATE INDEX i\_o\_clerk\_date ON orders(o\_clerk, o\_orderdate);**

**SELECT \* FROM orders  
WHERE o\_orderdate BETWEEN '1997-05-01' AND '1997-05-31'  
AND o\_clerk = 'Clerk#000001866';**

id	select type	table	type	possible keys	key	key len	ref	rows	extra
1	SIMPLE	orders	range	i_o_orderdate, i_o_clerk, i_o_clerk_date	i_o_clerk_date	20	NULL	14	Using index condition

```
mysql> SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' AND ...
...
15 rows in set (0.00 sec)
```

# Performance Schema: Query History

**UPDATE performance\_schema.setup\_consumers  
SET enabled='YES' WHERE name = 'events\_statements\_history';**

MySQL 5.7:  
Enabled by default

```
mysql> SELECT sql_text, (timer_wait)/1000000000.0 "t (ms)", rows_examined rows
        FROM performance_schema.events_statements_history ORDER BY timer_start;
```

sql_text	t (ms)	rows
SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ...	8.1690	1505
SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ...	7.2120	1505
SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ...	8.1613	1505
SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ...	7.0535	1505
CREATE INDEX i_o_clerk_date ON orders(o_clerk,o_orderdate)	82036.4190	0
SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ...	0.7259	15
SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ...	0.5791	15
SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ...	0.5423	15
SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ...	0.6031	15
SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ...	0.2710	15

# Program Agenda

- 1 ➤ Cost-based query optimization in MySQL
- 2 ➤ Tools for monitoring, analyzing, and tuning queries
- 3 ➤ Data access and index selection
- 4 ➤ Join optimizer**
- 5 ➤ Sorting
- 6 ➤ Influencing the optimizer

# Join Optimizer

## "Greedy search strategy"

- Goal: Given a JOIN of N tables, find the best JOIN ordering

$N!$  possible plans

- Strategy:

- Start with all 1-table plans (Sorted based on *size* and *key dependency*)

- Expand each plan with remaining tables

- Depth-first


- If “cost of partial plan” > “cost of best plan”:

- “prune” plan

- Heuristic pruning:

- Prune less promising partial plans

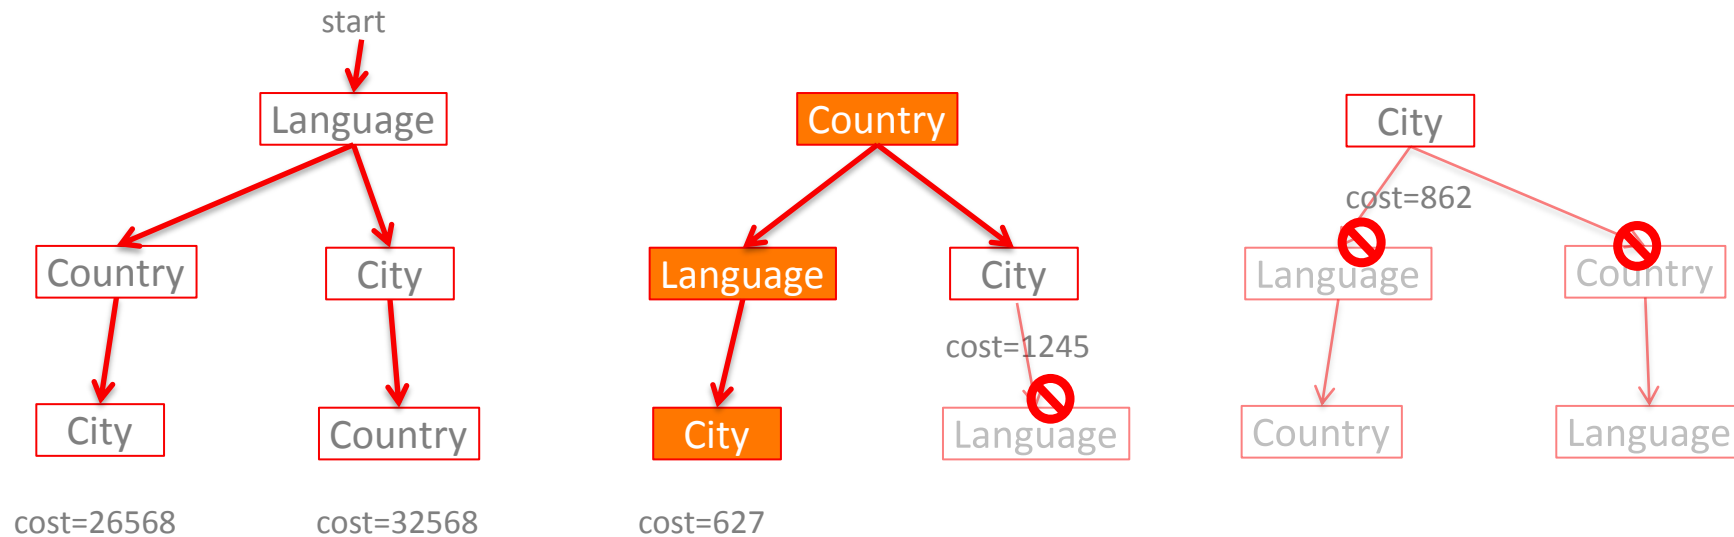
- May in rare cases miss most optimal plan (turn off with **set optimizer\_prune\_level = 0**)





# JOIN Optimizer Illustrated

```
SELECT City.Name, Language FROM Language, Country, City
WHERE City.CountryCode = Country.Code
      AND City.ID = Country.Capital
      AND City.Population >= 1000000
      AND Language.Country = Country.Code;
```



# Join Optimizer: Case study

## DBT-3 Query 8: National Market Share Query

```

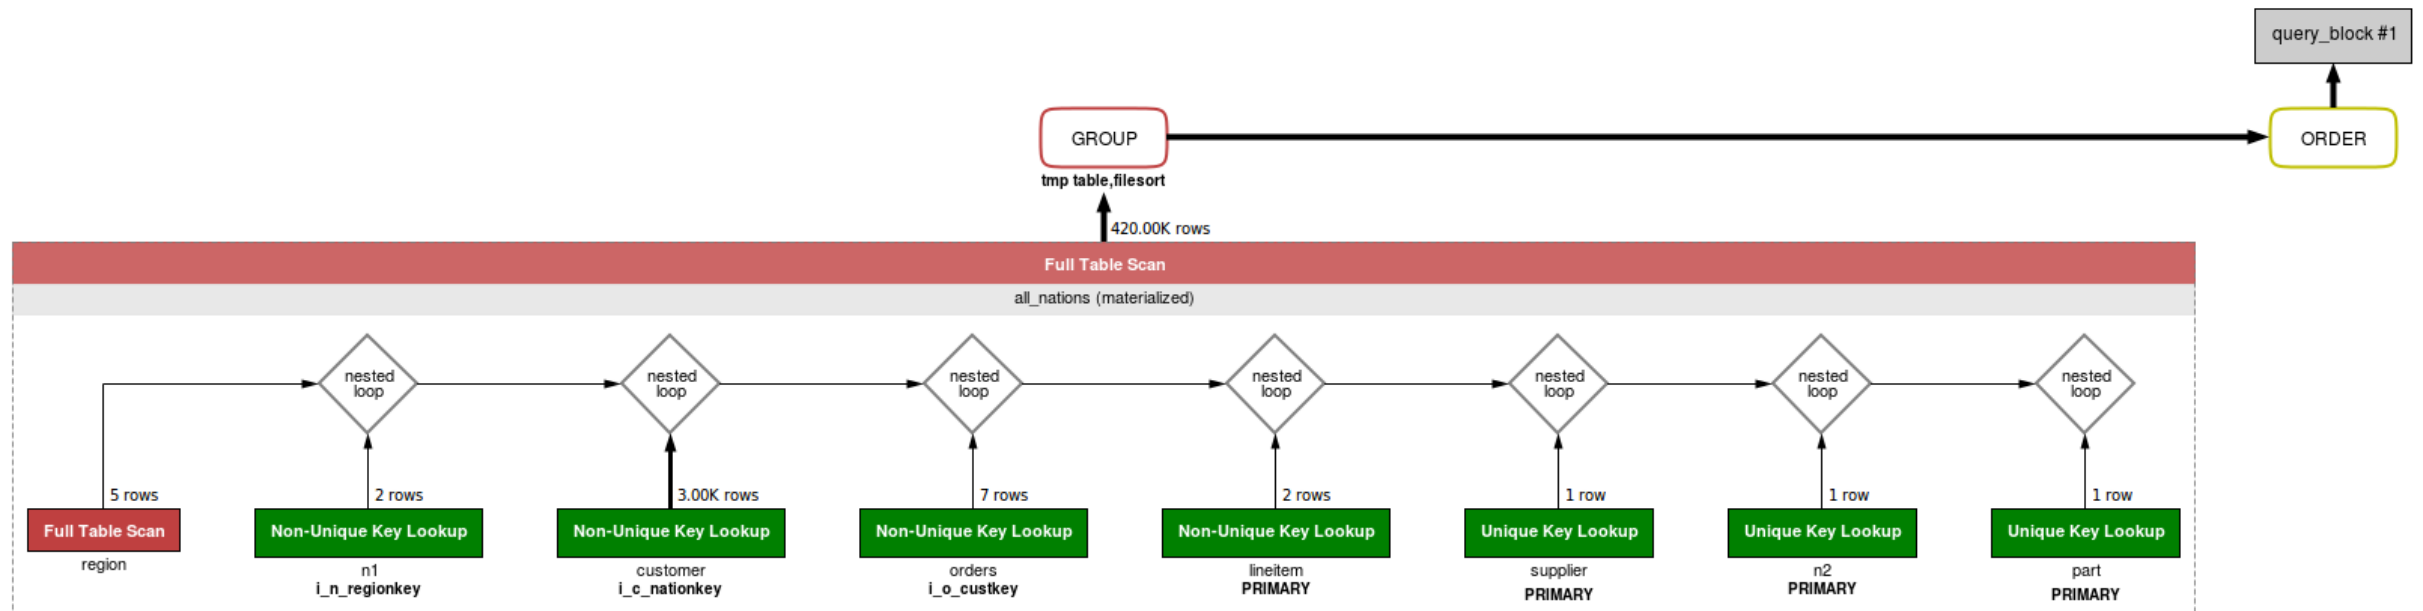
SELECT o_year, SUM(CASE WHEN nation = 'FRANCE' THEN volume ELSE 0 END) / SUM(volume) AS
    mkt_share
FROM (
    SELECT EXTRACT(YEAR FROM o_orderdate) AS o_year,
           l_extendedprice * (1 - l_discount) AS volume, n2.n_name AS nation
    FROM part
    JOIN lineitem ON p_partkey = l_partkey
    JOIN supplier ON s_suppkey = l_suppkey
    JOIN orders ON l_orderkey = o_orderkey
    JOIN customer ON o_custkey = c_custkey
    JOIN nation n1 ON c_nationkey = n1.n_nationkey
    JOIN region ON n1.n_regionkey = r_regionkey
    JOIN nation n2 ON s_nationkey = n2.n_nationkey
    WHERE r_name = 'EUROPE' AND o_orderdate BETWEEN '1995-01-01' AND '1996-12-31'
           AND p_type = 'PROMO BRUSHED STEEL'
) AS all_nations GROUP BY o_year ORDER BY o_year;

```

# Join Optimizer: Case Study

## MySQL Workbench: Visual EXPLAIN

Execution time: 21 seconds



# Join Optimizer: Case study

Force early processing of high selectivity predicates

```
SELECT o_year, SUM(CASE WHEN nation = 'FRANCE' THEN volume ELSE 0 END) / SUM(volume) AS
      mkt_share
```

```
FROM (
```

```
  SELECT EXTRACT(YEAR FROM o_orderdate) AS o_year,
         l_extendedprice * (1 - l_discount) AS volume, n2.n_name AS nation
  FROM part
```

part before lineitem

```
    STRAIGHT_JOIN lineitem ON p_partkey = l_partkey
    JOIN supplier ON s_suppkey = l_suppkey
    JOIN orders ON l_orderkey = o_orderkey
    JOIN customer ON o_custkey = c_custkey
    JOIN nation n1 ON c_nationkey = n1.n_nationkey
    JOIN region ON n1.n_regionkey = r_regionkey
    JOIN nation n2 ON s_nationkey = n2.n_nationkey
  WHERE r_name = 'EUROPE' AND o_orderdate BETWEEN '1995-01-01' AND '1996-12-31'
         AND p_type = 'PROMO BRUSHED STEEL'
```

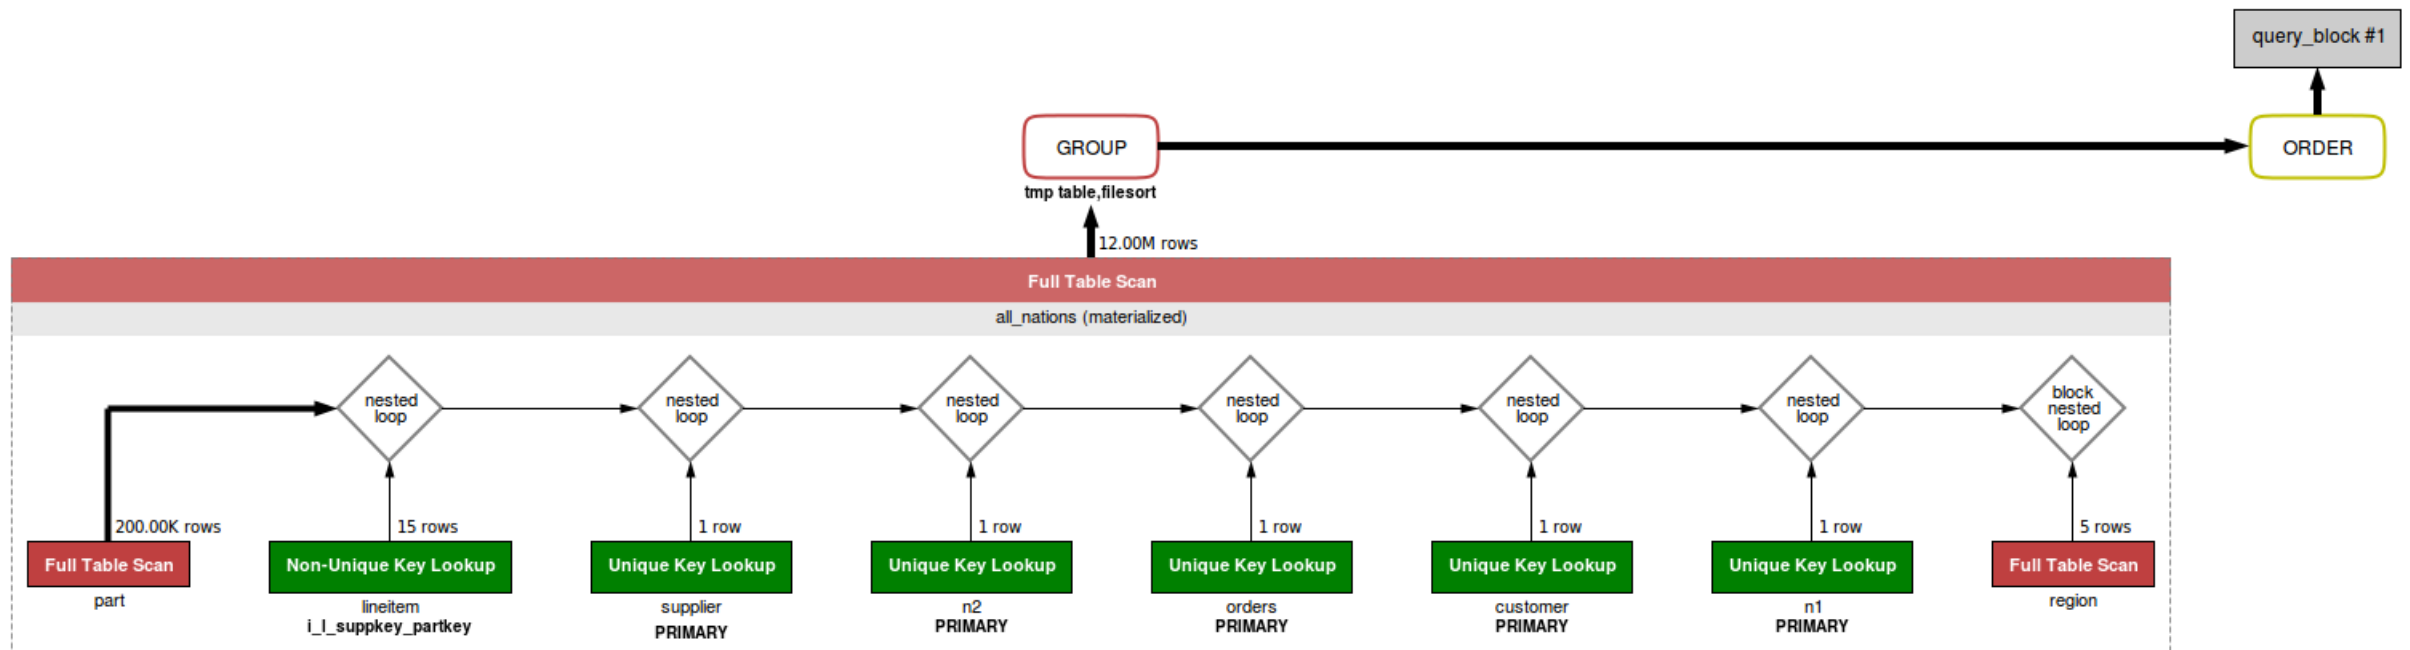
```
) AS all_nations GROUP BY o_year ORDER BY o_year;
```

Highest selectivity

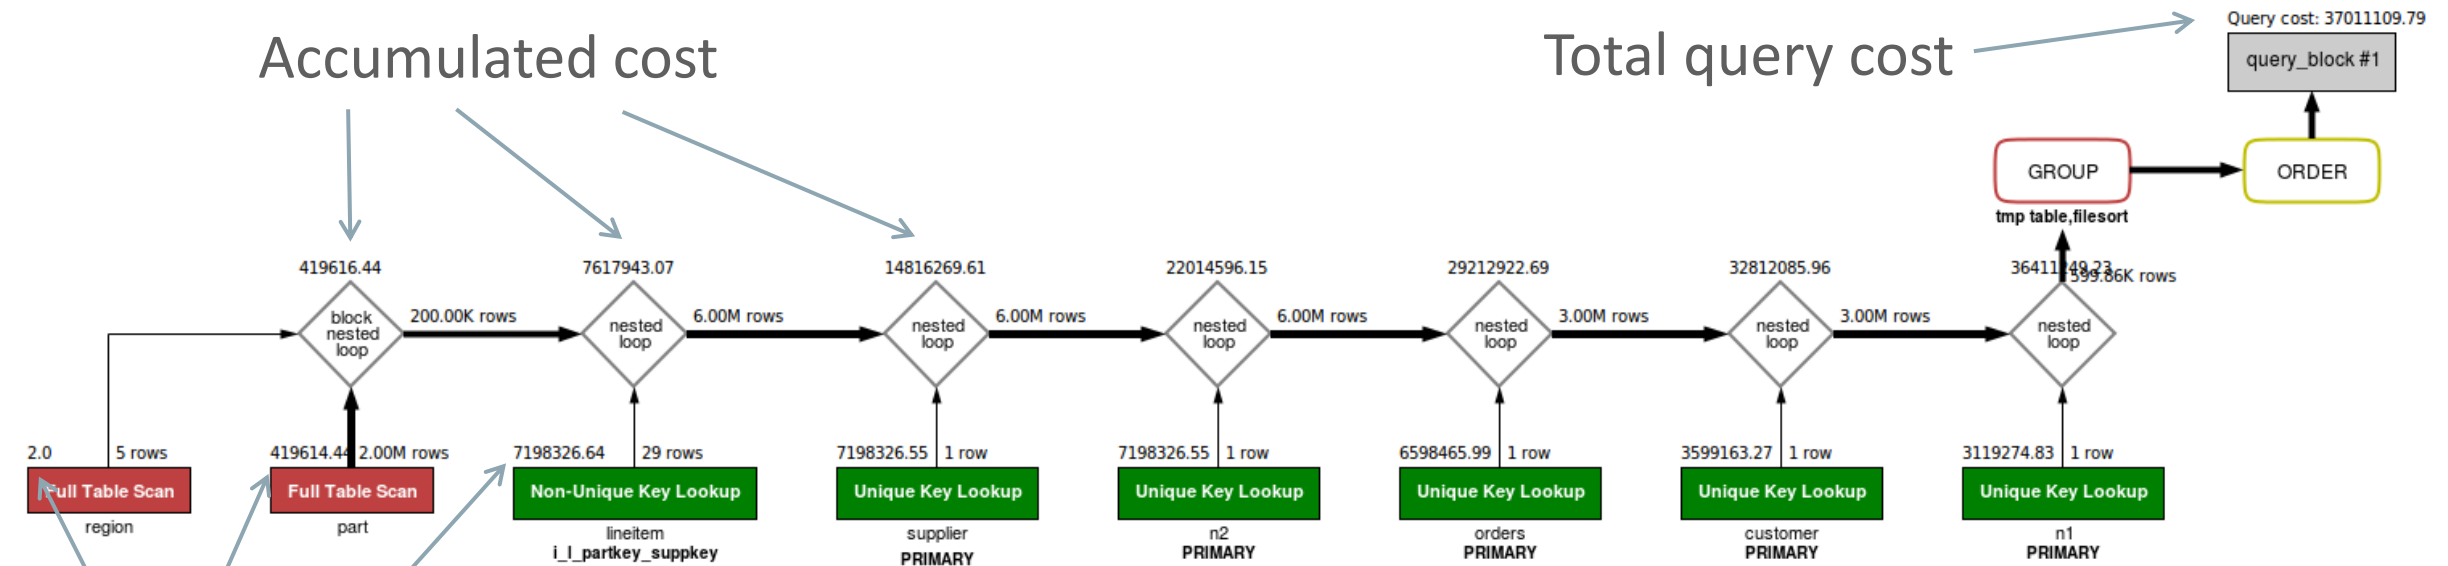
# Join Optimizer: Case study

## Improved join order

Execution time: 3 seconds



# MySQL 5.7: Cost Information in Structured EXPLAIN



Cost per table

## Improvements to Query 8 in MySQL 5.7:

- Filtering on non-indexed columns are taken into account
  - No need for hint to force **part** table to be processed early
- Merge derived tables into outer query
  - No temporary table

# Program Agenda

- 1 ➤ Cost-based query optimization in MySQL
- 2 ➤ Tools for monitoring, analyzing, and tuning queries
- 3 ➤ Data access and index selection
- 4 ➤ Join optimizer
- 5 ➤ **Sorting**
- 6 ➤ Influencing the optimizer

# ORDER BY Optimizations

- General solution; “Filesort”:
  - Store query result in temporary table before sorting
  - If data volume is large, may need to sort in several passes with intermediate storage on disk.
- Optimizations:
  - Take advantage of index to generate query result in sorted order
  - For “LIMIT  $n$ ” queries, maintain priority queue of  $n$  top items in memory instead of filesort. (MySQL 5.6)



# Filesort

```
SELECT * FROM orders ORDER BY o_totalprice ;
```

id	select type	table	type	possible keys	key	key len	ref	rows	extra
1	SIMPLE	orders	ALL	NULL	NULL	NULL	NULL	15000000	Using filesort

```
SELECT c_name, o_orderkey, o_totalprice
FROM orders JOIN customer ON c_custkey = o_custkey
WHERE c_acctbal < -1000 ORDER BY o_totalprice ;
```

id	select type	table	type	possible keys	key	key len	ref	rows	extra
1	SIMPLE	customer	ALL	PRIMARY	NULL	NULL	NULL	1500000	Using where; Using temporary; Using filesort
1	SIMPLE	orders	ref	i_o_custkey	i_o_custkey	5	...	7	NULL

# Filesort

## Status variables

Status variables related to sorting:

```
mysql> show status like 'Sort%';
```

Variable_name	Value
Sort_merge_passes	1
Sort_range	0
Sort_rows	136170
Sort_scan	1

>0: Intermediate storage on disk.

Consider increasing `sort_buffer_size`

Number of sort operations

(range scan or table/index scans)

Number of rows sorted

# Filesort

## Performance Schema

Sorting status per statement available from Performance Schema

```
mysql> SELECT sql_text,sort_merge_passes,sort_range,sort_rows,sort_scan
        FROM performance_schema.events_statements_history
        ORDER BY timer_start DESC LIMIT 1;
```

sql_text	sort_merge_passes	sort_range	sort_rows	sort_scan
SELECT ...	1	0	136170	1

# Filesort: Case Study

```
mysql> FLUSH STATUS;

Query OK, 0 rows affected (0.00 sec)

mysql> SELECT AVG(o_totalprice) FROM
  ( SELECT * FROM orders
    ORDER BY o_totalprice DESC
    LIMIT 100000) td;
+-----+
| AVG(o_totalprice) |
+-----+
| 398185.986158      |
+-----+
1 row in set (24.65 sec)
```

Unnecessary large data volume!

```
mysql> SHOW STATUS LIKE 'sort%';
+-----+-----+
| Variable_name | Value |
+-----+-----+
| Sort_merge_passes | 1432 |
| Sort_range       | 0     |
| Sort_rows        | 100000 |
| Sort_scan        | 1     |
+-----+-----+
4 rows in set (0.00 sec)
```

Many intermediate sorting steps!

# Filesort: Case Study

## Reduce amount of data to be sorted

```
mysql> SELECT AVG(o_totalprice) FROM (SELECT o_totalprice FROM orders ORDER BY
      o_totalprice DESC LIMIT 100000) td;
```

```
+-----+
| AVG(o_totalprice) |
+-----+
| 398185.986158      |
+-----+
1 row in set (8.18 sec)
```

```
mysql> SELECT sql_text, sort_merge_passes FROM performance_schema.
      events_statements_history ORDER BY timer_start DESC LIMIT 1;
```

```
+-----+-----+
| sql_text                                     | sort_merge_passes |
+-----+-----+
| SELECT AVG(o_totalprice) FROM (SELECT o_totalprice | 229 |
+-----+-----+
```

# Filesort: Case Study

Default is 256 kB

Increase sort buffer (1 MB)

```
mysql> SET sort_buffer_size = 1024*1024;
```

```
mysql> SELECT AVG(o_totalprice) FROM (SELECT o_totalprice FROM orders ORDER BY
o_totalprice DESC LIMIT 100000) td;
```

```
+-----+
| AVG(o_totalprice) |
+-----+
| 398185.986158      |
+-----+
```

1 row in set (7.24 sec)

```
mysql> SELECT sql_text, sort_merge_passes FROM performance_schema.
events_statements_history ORDER BY timer_start DESC LIMIT 1;
```

```
+-----+-----+
| sql_text                                     | sort_merge_passes |
+-----+-----+
| SELECT AVG(o_totalprice) FROM (SELECT o_totalprice | 57 |
+-----+-----+
```

# Filesort: Case Study

## Increase sort buffer even more (8 MB)

```
mysql> SET sort_buffer_size = 8*1024*1024;
```

```
mysql> SELECT AVG(o_totalprice) FROM (SELECT o_totalprice FROM orders ORDER BY
o_totalprice DESC LIMIT 100000) td;
```

```
+-----+
| AVG(o_totalprice) |
+-----+
| 398185.986158      |
+-----+
```

1 row in set (6.30 sec)

```
mysql> SELECT sql_text, sort_merge_passes FROM performance_schema.
events_statements_history ORDER BY timer_start DESC LIMIT 1;
```

```
+-----+-----+
| sql_text                                     | sort_merge_passes |
+-----+-----+
| SELECT AVG(o_totalprice) FROM (SELECT o_totalprice | 0 |
+-----+-----+
```

# Using Index to Avoid Sorting

```
CREATE INDEX i_o_totalprice ON orders(o_totalprice);
```

```
SELECT o_orderkey, o_totalprice FROM orders ORDER BY o_totalprice ;
```

id	select type	table	type	possible keys	key	key len	ref	rows	extra
1	SIMPLE	orders	index	NULL	i_o_totalprice	6	NULL	15000000	Using index

However, still (due to total cost):

```
SELECT * FROM orders ORDER BY o_totalprice ;
```

id	select type	table	type	possible keys	key	key len	ref	rows	extra
1	SIMPLE	orders	ALL	NULL	NULL	NULL	NULL	15000000	Using filesort



# Using Index to Avoid Sorting

## Case study revisited

```
SELECT AVG(o_totalprice) FROM
  (SELECT o_totalprice FROM orders ORDER BY o_totalprice DESC LIMIT 100000) td;
```

id	select type	table	Type	possible keys	key	key len	ref	rows	extra
1	PRIMARY	<derived2>	ALL	NULL	NULL	NULL	NULL	100000	NULL
2	DERIVED	orders	index	NULL	i_o_totalprice	6	NULL	15000000	Using index

```
mysql> SELECT AVG(o_totalprice) FROM (
  SELECT o_totalprice FROM orders
  ORDER BY o_totalprice DESC LIMIT 100000) td;
```

...

1 row in set (0.06 sec)

# Program Agenda

- 1 ➤ Cost-based query optimization in MySQL
- 2 ➤ Tools for monitoring, analyzing, and tuning queries
- 3 ➤ Data access and index selection
- 4 ➤ Join optimizer
- 5 ➤ Sorting
- 6 ➤ Influencing the optimizer

# Influencing the Optimizer

When the optimizer does not do what you want

- Add indexes
- Force use of specific indexes:
  - USE INDEX, FORCE INDEX, IGNORE INDEX
- Force specific join order:
  - STRAIGHT\_JOIN
- Adjust session variables
  - optimizer\_switch flags: set optimizer\_switch="index\_merge=off"
  - Buffer sizes: set sort\_buffer=8\*1024\*1024;
  - Other variables: set optimizer\_search\_depth = 10;

# MySQL 5.7: New Optimizer Hints

- New hint syntax:
  - **SELECT /\*+ HINT1(*args*) HINT2(*args*) \*/ ... FROM ...**
- New hints:
  - BKA(*tables*)/NO\_BKA(*tables*), BNL(*tables*)/NO\_BNL(*tables*)
  - MRR(*table indexes*)/NO\_MRR(*table indexes*)
  - SEMIJOIN/NO\_SEMIJOIN(*strategies*), SUBQUERY(*strategy*)
  - NO\_ICP(*table indexes*)
  - NO\_RANGE\_OPTIMIZATION(*table indexes*)
  - QB\_NAME(*name*)
- Finer granularity than **optimizer\_switch** session variable

# MySQL 5.7: Hint Example: SEMIJOIN

- No hint, optimizer chooses semi-join algorithm LooseScan:

**EXPLAIN SELECT \* FROM t2 WHERE t2.a IN (SELECT a FROM t3);**

id	select type	table	type	possible keys	key	key len	ref	rows	extra
1	SIMPLE	t3	index	a	a	4	NULL	3	Using where; LooseScan
1	SIMPLE	t2	ref	a	a	4	test.t3.a	1	Using index

- Disable semi-join with hint:

**EXPLAIN SELECT \* FROM t2 WHERE t2.a IN (SELECT /\*+ NO\_SEMIJOIN() \*/ a FROM t3);**

id	select type	table	type	possible keys	key	key len	ref	rows	extra
1	PRIMARY	t2	index	null	a	4	NULL	4	Using where; Using index
2	DEPENDENT SUBQUERY	t3	Index_subquery	a	a	4	func	1	Using index

# MySQL 5.7: Hint Example: SEMIJOIN

- Force Semi-join Materialization to be used

```
EXPLAIN SELECT /*+ SEMIJOIN(@subq MATERIALIZATION) */ * FROM t2
WHERE t2.a IN (SELECT /*+ QB_NAME(subq) */ a FROM t3);
```

id	select type	table	type	possible keys	key	key len	ref	rows	extra
1	SIMPLE	t2	index	a	a	4	NULL	4	Using where; Using index
1	SIMPLE	<subquery2>	eq_ref	<auto_key>	<auto_key>	4	test.t2.a	1	NULL
2	MATERIALIZED	t3	index	a	a	4	NULL	3	Using index

# MySQL 5.7: Query Rewrite Plugin

- Rewrite problematic queries without the need to make application changes
  - Add hints
  - Modify join order
  - Much more ...

- Add rewrite rules to table:

```
INSERT INTO query_rewrite.rewrite_rules (pattern, replacement ) VALUES  
("SELECT * FROM t1 WHERE a > ? AND b = ?",  
 "SELECT * FROM t1 FORCE INDEX (a_idx) WHERE a > ? AND b = ?");
```

- New pre- and post-parse query rewrite APIs
  - Users can write their own plug-ins

# More information

- My blog:
  - <http://oysteing.blogspot.com/>
- Optimizer team blog:
  - <http://mysqloptimizerteam.blogspot.com/>
- MySQL Server Team blog
  - <http://mysqlserverteam.com/>
- MySQL forums:
  - Optimizer & Parser: <http://forums.mysql.com/list.php?115>
  - Performance: <http://forums.mysql.com/list.php?24>



Q+A



**Thank You for Joining Us Today  
Please Move to the Next Session.**

# Integrated Cloud

## Applications & Platform Services

ORACLE®