



Virtual Developer Day—MySQL
Brought to You by Oracle Technology Network

ORACLE

Profiling with Performance Schema

Mark Leith
Software Development Senior Manager

Program Agenda

- An Introduction to Performance Schema
- Performance Schema Configuration
- Profiling General Instance Activity
- Profiling Statement Activity

THE FOLLOWING IS INTENDED TO OUTLINE OUR GENERAL PRODUCT DIRECTION. IT IS INTENDED FOR INFORMATION PURPOSES ONLY, AND MAY NOT BE INCORPORATED INTO ANY CONTRACT. IT IS NOT A COMMITMENT TO DELIVER ANY MATERIAL, CODE, OR FUNCTIONALITY, AND SHOULD NOT BE RELIED UPON IN MAKING PURCHASING DECISION. THE DEVELOPMENT, RELEASE, AND TIMING OF ANY FEATURES OR FUNCTIONALITY DESCRIBED FOR ORACLE'S PRODUCTS REMAINS AT THE SOLE DISCRETION OF ORACLE.

An Introduction to Performance Schema



An Introduction to Performance Schema

- Performance Schema Overview
- A new default ***performance_schema*** schema within MySQL
- Tables use new ***PERFORMANCE_SCHEMA*** storage engine
 - Real engine, unlike INFORMATION_SCHEMA
- Records various run time statistics via in-built instrumentation points
- All recorded statistics are stored in fixed size ring buffers in memory
- Most instrumentation configuration can be done dynamically

An Introduction to Performance Schema

- Instrumentation Points
- At it's core, Performance Schema tracks **latency** for various **events**
- Latency is exposed to **picosecond** precision (a *trillionth* of a second)
- In 5.5:
 - File I/O, Mutexes, Read/Write Locks, Conditions
- In 5.6:
 - Network I/O, Table I/O, Table Locks, Stages, Statements, Idle
- Also tracks other data as appropriate, like bytes, source position, etc.

An Introduction to Performance Schema

- Wait Events

▪ Table Locks	wait/lock/table/%	(5.6)
▪ Network IO	wait/io/socket/%	(5.6)
▪ Table IO	wait/io/table/%	(5.6)
▪ File IO	wait/io/file/%	(5.5)
▪ Mutexes	wait/synch/mutex/%	(5.5)
▪ Read/Write Locks	wait/synch/rwlock/%	(5.5)
▪ Conditions	wait/synch/cond/%	(5.5)

An Introduction to Performance Schema

▪ Raw Wait Events – A Table IO Event Example

```
mysql> select * from events_waits_history where event_name like 'wait/io/table%'\G
***** 1. row *****
      THREAD_ID: 352950
      EVENT_ID: 913915 (Event context, who, and in what level)
      END_EVENT_ID: 913915
      EVENT_NAME: wait/io/table/sql/handler (Event type & origination)
      SOURCE: handler.cc:2716
      TIMER_START: 1301573336396109228
      TIMER_END: 1301573336396938344 (Timing – 829 nanoseconds)
      TIMER_WAIT: 829116
      SPINS: NULL
      OBJECT_SCHEMA: mem30__quan (Database object the event was against)
      OBJECT_NAME: normalized_statements_by_server_by_schema_data
      INDEX_NAME: PRIMARY
      OBJECT_TYPE: TABLE
      OBJECT_INSTANCE_BEGIN: 360798712 (Object type, and individual instance)
      NESTING_EVENT_ID: 484888
      NESTING_EVENT_TYPE: STAGE (Whether the event was nested within another)
      OPERATION: fetch
      NUMBER_OF_BYTES: NULL (Further info on the event, it's type, size, flags)
      FLAGS: NULL
```

An Introduction to Performance Schema

- Statement and Stage Events (5.6+)
- Statements have two types of event
 - SQL Statements `statement/sql/%`
 - COM Commands `statement/com/%`
- Stages are the thread states (like SHOW PROFILE, but available across connections)
 - One form so far `stage/sql/%`
 - Apart from... `stage/mysys/Waiting for table level lock`

An Introduction to Performance Schema

- Raw Stage Events

```
mysql> select * from events_stages_history\G
***** 1. row *****
      THREAD_ID: 378012
      EVENT_ID: 273
      END_EVENT_ID: 273
      EVENT_NAME: stage/sql/preparing
      SOURCE: sql_optimizer.cc:483
      TIMER_START: 1380151756100562214
      TIMER_END: 1380151756107845688
      TIMER_WAIT: 7283474
      NESTING_EVENT_ID: 249
      NESTING_EVENT_TYPE: STATEMENT
***** 2. row *****
      THREAD_ID: 378012
      EVENT_ID: 274
      END_EVENT_ID: 274
      EVENT_NAME: stage/sql/executing
      SOURCE: sql_executor.cc:112
      TIMER_START: 1380151756107845688
      TIMER_END: 1380151756108507822
      TIMER_WAIT: 662134
      NESTING_EVENT_ID: 249
      NESTING_EVENT_TYPE: STATEMENT
```

An Introduction to Performance Schema

▪ Raw Statement Events

```
mysql> select * from events_statements_history limit 1\G
***** 1. row *****
      THREAD_ID: 378012
        EVENT_ID: 249
    END_EVENT_ID: 288
    EVENT_NAME: statement/sql/select
        SOURCE: mysqld.cc:907
    TIMER_START: 1380151755620846131
    TIMER_END: 1380151756820301872
    TIMER_WAIT: 1199455741
      LOCK_TIME: 0
    SQL_TEXT: select * from events_stages_history
      DIGEST: f13b9e15a70df2108f6aa343860734fa
    DIGEST_TEXT: SELECT * FROM events_stages_history
    CURRENT_SCHEMA: performance_schema
      OBJECT_TYPE: NULL
    OBJECT_SCHEMA: NULL
      OBJECT_NAME: NULL
OBJECT_INSTANCE_BEGIN: NULL
      MYSQL_ERRNO: 0
    RETURNED_SQLSTATE: NULL
      MESSAGE_TEXT: NULL
```

```
      ERRORS: 0
    WARNINGS: 0
    ROWS_AFFECTED: 0
      ROWS_SENT: 125
    ROWS_EXAMINED: 125
CREATED_TMP_DISK_TABLES: 0
    CREATED_TMP_TABLES: 0
      SELECT_FULL_JOIN: 0
    SELECT_FULL_RANGE_JOIN: 0
      SELECT_RANGE: 0
    SELECT_RANGE_CHECK: 0
      SELECT_SCAN: 1
    SORT_MERGE_PASSES: 0
      SORT_RANGE: 0
      SORT_ROWS: 0
      SORT_SCAN: 0
      NO_INDEX_USED: 1
    NO_GOOD_INDEX_USED: 0
      NESTING_EVENT_ID: NULL
    NESTING_EVENT_TYPE: NULL
```

An Introduction to Performance Schema

■ Setup Tables

- Used to define certain configuration dynamically
- Can perform DML against these tables

```
mysql> SELECT table_name
-> FROM information_schema.tables
-> WHERE table_schema = 'performance_schema'
-> AND table_name LIKE 'setup%';
```

table_name
setup_actors
setup_consumers
setup_instruments
setup_objects
setup_timers

An Introduction to Performance Schema

Raw Data Tables

- Expose events, objects, or instances of instruments in a raw manner
- Allow seeing a (brief) history of raw event metrics as well

```
mysql> SELECT table_name
-> FROM information_schema.tables
-> WHERE table_schema = 'performance_schema'
-> AND table_name NOT LIKE 'setup%'
-> AND table_name NOT LIKE '%summary%';
```

table_name
accounts
cond_instances
events_stages_current
events_stages_history
events_stages_history_long
events_statements_current
events_statements_history
events_statements_history_long
events_waits_current
events_waits_history
events_waits_history_long
file_instances
host_cache
hosts
mutex_instances
performance_timers
rwlock_instances
session_account_connect_attrs
session_connect_attrs
socket_instances
threads
users

An Introduction to Performance Schema

Summary Tables

- Summarize event information over various dimensions
- Useful for longer term monitoring of activity

```
mysql> SELECT table_name
-> FROM information_schema.tables
-> WHERE table_schema = 'performance_schema'
-> AND table_name LIKE '%summary%';
```

table_name
events_stages_summary_by_account_by_event_name
events_stages_summary_by_host_by_event_name
events_stages_summary_by_thread_by_event_name
events_stages_summary_by_user_by_event_name
events_stages_summary_global_by_event_name
events_statements_summary_by_account_by_event_name
events_statements_summary_by_digest
events_statements_summary_by_host_by_event_name
events_statements_summary_by_thread_by_event_name
events_statements_summary_by_user_by_event_name
events_statements_summary_global_by_event_name
events_waits_summary_by_account_by_event_name
events_waits_summary_by_host_by_event_name
events_waits_summary_by_instance
events_waits_summary_by_thread_by_event_name
events_waits_summary_by_user_by_event_name
events_waits_summary_global_by_event_name
file_summary_by_event_name
file_summary_by_instance
objects_summary_global_by_type
socket_summary_by_event_name
socket_summary_by_instance
table_io_waits_summary_by_index_usage
table_io_waits_summary_by_table
table_lock_waits_summary_by_table

Performance Schema Configuration



Performance Schema Configuration

- Configuration Options
- There are three distinct types of configuration
- A way to configure ***how much memory to allocate*** for instruments
 - Set with various system variables in options files, not dynamic
- A way to ***enable/disable instrumentation at startup***
 - Again set within options files, only available as of 5.6
- A way to ***enable/disable instrumentation dynamically***
 - Set by updating various “setup” tables dynamically

Performance Schema Configuration

- Memory Configuration
 - Memory usage falls in to two distinct categories
 - How many “classes” and “instances” of event types to track
 - Affects all data available in performance schema
 - How much summary and history data you want to track
 - Affects how long you can see data for

Performance Schema Configuration

■ Class / Instance Config

- Classes count instrument implementations -
“wait/io/file/sql/binlog”
- Instances are the different runtime instances of them -
“/data/mysql/binlog.000001”

```
mysql> SELECT variable_name, variable_value  
-> FROM information_schema.global_variables  
-> WHERE variable_name LIKE 'perf%classes'  
-> OR variable_name LIKE 'perf%instances'  
-> ORDER BY variable_name;
```

variable_name	variable_value
PERFORMANCE_SCHEMA_MAX_COND_CLASSES	80
PERFORMANCE_SCHEMA_MAX_COND_INSTANCES	723
PERFORMANCE_SCHEMA_MAX_FILE_CLASSES	50
PERFORMANCE_SCHEMA_MAX_FILE_INSTANCES	1556
PERFORMANCE_SCHEMA_MAX_MUTEX_CLASSES	200
PERFORMANCE_SCHEMA_MAX_MUTEX_INSTANCES	3112
PERFORMANCE_SCHEMA_MAX_RWLOCK_CLASSES	30
PERFORMANCE_SCHEMA_MAX_RWLOCK_INSTANCES	1667
PERFORMANCE_SCHEMA_MAX_SOCKET_CLASSES	10
PERFORMANCE_SCHEMA_MAX_SOCKET_INSTANCES	123
PERFORMANCE_SCHEMA_MAX_STAGE_CLASSES	150
PERFORMANCE_SCHEMA_MAX_STATEMENT_CLASSES	169
PERFORMANCE_SCHEMA_MAX_TABLE_INSTANCES	445
PERFORMANCE_SCHEMA_MAX_THREAD_CLASSES	50
PERFORMANCE_SCHEMA_MAX_THREAD_INSTANCES	167

Performance Schema Configuration

■ Class / Instance Config

- Monitor the ***_lost** status variables opposite to see whether these need tweaking
- You will generally not need to modify classes – but instances could need tuning

```
mysql> SELECT variable_name, variable_value  
-> FROM information_schema.global_status  
-> WHERE variable_name LIKE 'perf%classes_lost'  
-> OR variable_name LIKE 'perf%instances_lost';
```

variable_name	variable_value
PERFORMANCE_SCHEMA_COND_CLASSES_LOST	0
PERFORMANCE_SCHEMA_COND_INSTANCES_LOST	0
PERFORMANCE_SCHEMA_FILE_CLASSES_LOST	0
PERFORMANCE_SCHEMA_FILE_INSTANCES_LOST	438056
PERFORMANCE_SCHEMA_MUTEX_CLASSES_LOST	0
PERFORMANCE_SCHEMA_MUTEX_INSTANCES_LOST	0
PERFORMANCE_SCHEMA_RWLOCK_CLASSES_LOST	0
PERFORMANCE_SCHEMA_RWLOCK_INSTANCES_LOST	0
PERFORMANCE_SCHEMA_SOCKET_CLASSES_LOST	0
PERFORMANCE_SCHEMA_SOCKET_INSTANCES_LOST	0
PERFORMANCE_SCHEMA_STAGE_CLASSES_LOST	0
PERFORMANCE_SCHEMA_STATEMENT_CLASSES_LOST	0
PERFORMANCE_SCHEMA_TABLE_INSTANCES_LOST	0
PERFORMANCE_SCHEMA_THREAD_CLASSES_LOST	0
PERFORMANCE_SCHEMA_THREAD_INSTANCES_LOST	0

Performance Schema Configuration

History / Summary Sizes

- Define total rows per type
- The *history_size variables are per-thread
- The *history_long_size is total rows
- Only settable in my.cnf/ini

```
mysql> SELECT variable_name, variable_value  
-> FROM information_schema.global_variables  
-> WHERE variable_name LIKE 'perf%'  
-> AND variable_name NOT LIKE '%instances'  
-> AND variable_name NOT LIKE '%classes'  
-> ORDER BY variable_name;
```

variable_name	variable_value
PERFORMANCE_SCHEMA	ON
PERFORMANCE_SCHEMA_ACCOUNTS_SIZE	100
PERFORMANCE_SCHEMA_DIGESTS_SIZE	10000
PERFORMANCE_SCHEMA_EVENTS_STAGES_HISTORY_LONG_SIZE	10000
PERFORMANCE_SCHEMA_EVENTS_STAGES_HISTORY_SIZE	10
PERFORMANCE_SCHEMA_EVENTS_STATEMENTS_HISTORY_LONG_SIZE	10000
PERFORMANCE_SCHEMA_EVENTS_STATEMENTS_HISTORY_SIZE	10
PERFORMANCE_SCHEMA_EVENTS_WAITS_HISTORY_LONG_SIZE	10000
PERFORMANCE_SCHEMA_EVENTS_WAITS_HISTORY_SIZE	10
PERFORMANCE_SCHEMA_HOSTS_SIZE	100
PERFORMANCE_SCHEMA_MAX_FILE_HANDLES	32768
PERFORMANCE_SCHEMA_MAX_TABLE_HANDLES	256
PERFORMANCE_SCHEMA_SESSION_CONNECT_ATTRS_SIZE	512
PERFORMANCE_SCHEMA_SETUP_ACTORS_SIZE	100
PERFORMANCE_SCHEMA_SETUP_OBJECTS_SIZE	100
PERFORMANCE_SCHEMA_USERS_SIZE	100

Performance Schema Configuration

■ setup_timers

- Define the granularity of timing
- Set different timers for different types of instrumentation

```
mysql> SELECT * FROM setup_timers;
+-----+-----+
| NAME      | TIMER_NAME |
+-----+-----+
| idle      | MICROSECOND |
| wait      | CYCLE      |
| stage     | MICROSECOND |
| statement | MICROSECOND |
+-----+-----+
4 rows in set (0.00 sec)
```

```
mysql> UPDATE setup_timers
-> SET timer_name = 'millisecond'
-> WHERE name = 'idle';
Query OK, 1 row affected (0.02 sec)
Rows matched: 1  Changed: 1  Warnings: 0
```

```
mysql> SELECT * FROM setup_timers;
+-----+-----+
| NAME      | TIMER_NAME |
+-----+-----+
| idle      | MILLISECOND |
| wait      | CYCLE      |
| stage     | MICROSECOND |
| statement | MICROSECOND |
+-----+-----+
4 rows in set (0.00 sec)
```

Performance Schema Configuration

performance_timers

- Show which timers are available
- Show the performance characteristics of the timer on the platform MySQL installed on

```
mysql> SELECT * FROM performance_timers\G
***** 1. row *****
      TIMER_NAME: CYCLE
      TIMER_FREQUENCY: 3090445149      (Good Choice)
      TIMER_RESOLUTION: 1
      TIMER_OVERHEAD: 21
***** 2. row *****
      TIMER_NAME: NANOSECOND
      TIMER_FREQUENCY: NULL      (Not Available)
      TIMER_RESOLUTION: NULL
      TIMER_OVERHEAD: NULL
***** 3. row *****
      TIMER_NAME: MICROSECOND
      TIMER_FREQUENCY: 3020537      (Good Choice)
      TIMER_RESOLUTION: 1
      TIMER_OVERHEAD: 21
***** 4. row *****
      TIMER_NAME: MILLISECOND
      TIMER_FREQUENCY: 1035      (High Overhead)
      TIMER_RESOLUTION: 15
      TIMER_OVERHEAD: 151
***** 5. row *****
      TIMER_NAME: TICK
      TIMER_FREQUENCY: 959      (Just... Don't.)
      TIMER_RESOLUTION: 1
      TIMER_OVERHEAD: 9223372036854775807
5 rows in set (0.00 sec)
```

Performance Schema Configuration

performance_timers

- Show which timers are available
- Show the performance characteristics of the timer on the platform MySQL installed on

```
mysql> SELECT * FROM performance_timers\G
***** 1. row *****
      TIMER_NAME: CYCLE
      TIMER_FREQUENCY: 3090445149      (Good Choice)
      TIMER_RESOLUTION: 1
      TIMER_OVERHEAD: 21
***** 2. row *****
      TIMER_NAME: NANOSECOND
      TIMER_FREQUENCY: NULL      (Not Available)
      TIMER_RESOLUTION: NULL
      TIMER_OVERHEAD: NULL
***** 3. row *****
      TIMER_NAME: MICROSECOND
      TIMER_FREQUENCY: 3020537      (Good Choice)
      TIMER_RESOLUTION: 1
      TIMER_OVERHEAD: 21
***** 4. row *****
      TIMER_NAME: MILLISECOND
      TIMER_FREQUENCY: 1035      (High Overhead)
      TIMER_RESOLUTION: 15
      TIMER_OVERHEAD: 151
***** 5. row *****
      TIMER_NAME: TICK
      TIMER_FREQUENCY: 959      (Just... Don't.)
      TIMER_RESOLUTION: 1
      TIMER_OVERHEAD: 9223372036854775807
5 rows in set (0.00 sec)
```

Performance Schema Configuration

■ setup_instruments

- Turn on/off individual instruments, and timing, dynamically
- Use UPDATE to modify

```
mysql> SELECT * FROM setup_instruments LIMIT 20;
```

NAME	ENABLED	TIMED
wait/synch/mutex/sql/PAGE::lock	YES	YES
wait/synch/mutex/sql/TC_LOG_MMAP::LOCK_sync	YES	YES
wait/synch/mutex/sql/TC_LOG_MMAP::LOCK_active	YES	YES
wait/synch/mutex/sql/TC_LOG_MMAP::LOCK_pool	YES	YES
wait/synch/mutex/sql/LOCK_des_key_file	YES	YES
wait/synch/mutex/sql/BINARY_LOG::LOCK_commit	YES	YES
wait/synch/mutex/sql/BINARY_LOG::LOCK_commit_queue	YES	YES
wait/synch/mutex/sql/BINARY_LOG::LOCK_done	YES	YES
wait/synch/mutex/sql/BINARY_LOG::LOCK_flush_queue	YES	YES
wait/synch/mutex/sql/BINARY_LOG::LOCK_index	YES	YES
wait/synch/mutex/sql/BINARY_LOG::LOCK_log	YES	YES
wait/synch/mutex/sql/BINARY_LOG::LOCK_sync	YES	YES
wait/synch/mutex/sql/BINARY_LOG::LOCK_sync_queue	YES	YES
wait/synch/mutex/sql/RELAY_LOG::LOCK_commit	YES	YES
wait/synch/mutex/sql/RELAY_LOG::LOCK_commit_queue	YES	YES
wait/synch/mutex/sql/RELAY_LOG::LOCK_done	YES	YES
wait/synch/mutex/sql/RELAY_LOG::LOCK_flush_queue	YES	YES
wait/synch/mutex/sql/RELAY_LOG::LOCK_index	YES	YES
wait/synch/mutex/sql/RELAY_LOG::LOCK_log	YES	YES
wait/synch/mutex/sql/RELAY_LOG::LOCK_sync	YES	YES

Performance Schema Configuration

■ setup_consumers

- Define how much to record in history (again, UPDATE to modify)
- **Statement Digests** also enabled

```
mysql> select * from setup_consumers;
```

NAME	ENABLED
events_stages_current	NO
events_stages_history	NO
events_stages_history_long	NO
events_statements_current	YES
events_statements_history	NO
events_statements_history_long	NO
events_waits_current	NO
events_waits_history	NO
events_waits_history_long	NO
global_instrumentation	YES
thread_instrumentation	YES
statements_digest	YES

Performance Schema Configuration

■ setup_actors (New in 5.6)

- Define exactly which connections to monitor
- By default monitors all connections (% is a wildcard)
- Modifications only apply to new connections

```
mysql> SELECT * FROM setup_actors;
+-----+-----+-----+
| HOST | USER | ROLE |
+-----+-----+-----+
| %    | %    | %    |
+-----+-----+-----+
1 row in set (0.00 sec)

mysql> DELETE FROM setup_actors;
Query OK, 1 row affected (0.00 sec)

mysql> INSERT INTO setup_actors
      -> VALUES ('%', 'mark', '%');
Query OK, 1 row affected (0.00 sec)

mysql> SELECT * FROM setup_actors;
+-----+-----+-----+
| HOST | USER | ROLE |
+-----+-----+-----+
| %    | mark | %    |
+-----+-----+-----+
1 row in set (0.00 sec)
```

Performance Schema Configuration

- **setup_objects (New in 5.6)**
- Define exactly which database objects to monitor
- Filters standard databases by default
- Allows explicit enable/disable, matches on most specific

```
mysql> SELECT * FROM setup_objects;
```

OBJECT_TYPE	OBJECT_SCHEMA	OBJECT_NAME	ENABLED	TIMED
TABLE	mysql	%	NO	NO
TABLE	performance_schema	%	NO	NO
TABLE	information_schema	%	NO	NO
TABLE	%	%	YES	YES

```
4 rows in set (0.00 sec)
```

```
mysql> INSERT INTO setup_objects  
-> VALUES ('TABLE', 'myschema', 'logging', 'NO', 'NO');  
Query OK, 1 row affected (0.00 sec)
```

```
mysql> SELECT * FROM setup_objects;
```

OBJECT_TYPE	OBJECT_SCHEMA	OBJECT_NAME	ENABLED	TIMED
TABLE	mysql	%	NO	NO
TABLE	performance_schema	%	NO	NO
TABLE	information_schema	%	NO	NO
TABLE	%	%	YES	YES
TABLE	myschema	logging	NO	NO

```
5 rows in set (0.00 sec)
```

Performance Schema Configuration

- Altering Enabled Instrumentation At Startup (5.6)
- Enable/Disable instruments at startup
 - `performance_schema_instrument='instrument_name=value'`
 - Value = [on | true | 1] / [off | false | 0] / [counted] (no timing)
 - Instruments can include wildcards “wait/synch/mutex/%”
- Enable/Disable consumers at startup
 - `performance_schema_consumer_consumer_name=value`
 - `performance_schema_consumer_events_stages_current=on`

Performance Schema Configuration

- Finally, grab `ps_helper`!
- A collection of functions / views / procedures to help with P_S
- I will be using the functions in the examples following for formatting
- Versions available for both 5.5 and 5.6
- http://www.markleith.co.uk/ps_helper/

Profiling General Instance Activity



Profiling General Instance Activity

- Defining what you want/need to monitor
- It's easy to just enable everything, but there are overhead concerns
 - On busy systems mutexes can be locked *millions of times a second*
- For busy systems, it's best to just pick higher latency event types
 - Statements, Stages, Table IO, File IO, maybe Network IO
 - It's good to enable all *current consumers, and statement/stage history
- For concurrency issues toggle other instruments on as needed basis

Profiling General Instance Activity

- Different Ways of Monitoring
- Once you've narrowed down what you are interested in, there are two ways to start monitoring
- View raw data in the summary views
 - This gives you an overall picture of usage on the instance
- Snapshot data, and compute deltas over time
 - This gives you an idea of the rates of changes for events
- Let's start with viewing raw summary data..

Profiling General Instance Activity

- Top Waits by Latency –
events_waits_summary_global_by_event_name

```
mysql> SELECT event_name AS event,  
->      count_star AS total_events,  
->      format_time(sum_timer_wait) AS total_latency,  
->      format_time(avg_timer_wait) AS avg_latency,  
->      format_time(max_timer_wait) AS max_latency  
-> FROM performance_schema.events_waits_summary_global_by_event_name  
-> WHERE event_name != 'idle'  
-> ORDER BY sum_timer_wait DESC LIMIT 5;
```

Only using InnoDB, this
must be from temp tables

event	total_events	total_latency	avg_latency	max_latency
wait/io/file/myisam/dfile	3623721056	00:47:49.09	791.70 ns	312.96 ms
wait/io/table/sql/handler	69879369	00:44:55.81	38.58 us	879.49 ms
wait/io/file/innodb/innodb_log_file	28286631	00:37:57.13	80.50 us	476.00 ms
wait/io/socket/sql/client_connection	201937035	00:18:44.59	5.57 us	1.27 s
wait/io/file/innodb/innodb_data_file	2835077	00:08:15.34	174.72 us	455.22 ms

Profiling General Instance Activity

- Analyzing Global Waits
- Some mutex events that can affect global throughput (if high in list):
 - wait/synch/mutex/innodb/buf_pool_mutex
 - ***Increase innodb_buffer_pool_instances***
 - wait/synch/mutex/sql/Query_cache::structure_guard_mutex
 - ***Disable the Query Cache***
 - wait/synch/mutex/myisam/MYISAM_SHARE::intern_lock
 - ***Use InnoDB ;)***

Profiling General Instance Activity

- Analyzing Global Waits
- Some File IO events to watch for:
 - wait/io/file/sql/FRM
 - *Tune table_open_cache / table_definition_cache*
 - wait/io/file/sql/file_parser (View definition parsing)
 - *If high on 5.5, upgrade to 5.6 (which caches these)*
 - wait/io/file/sql/query_log / wait/io/file/sql/slow_log
 - *Disable General / Slow query logs*

Profiling General Instance Activity

▪ Top Files by Total IO – file_summary_by_instance

```
mysql> SELECT format_path(file_name) AS file,  
-> count_read,  
-> format_bytes(sum_number_of_bytes_read) AS total_read,  
-> format_bytes(IFNULL(sum_number_of_bytes_read / count_read, 0)) AS avg_read,  
-> count_write,  
-> format_bytes(sum_number_of_bytes_write) AS total_written,  
-> format_bytes(IFNULL(sum_number_of_bytes_write / count_write, 0.00)) AS avg_write,  
-> format_bytes(sum_number_of_bytes_read + sum_number_of_bytes_write) AS total,  
-> IFNULL(ROUND(100-((sum_number_of_bytes_read/(sum_number_of_bytes_read+sum_number_of_bytes_write))*100), 2), 0.00) AS write_pct  
-> FROM performance_schema.file_summary_by_instance  
-> ORDER BY sum_number_of_bytes_read + sum_number_of_bytes_write DESC LIMIT 5;
```

High IO on per-tablespace tables
are candidates for a separate
mountpoint / disk
with “DATA DIRECTORY”

file	count_read	total_read	avg_read	count_write	total_written	avg_write	total	write_pct
@@datadir/ibdata1	801	14.48 MiB	18.52 KiB	46197285	1.66 TiB	38.66 KiB	1.66 TiB	100.00
@@datadir/mem_events/events.ibd	293	4.58 MiB	16.00 KiB	8066205	123.42 GiB	16.04 KiB	123.43 GiB	100.00
@@datadir/mem_qan/normalized_statements_by_server_by_schema_data.ibd	5	80.00 KiB	16.00 KiB	4536985	70.61 GiB	16.32 KiB	70.61 GiB	100.00
@@datadir/ib_logfile0	6	68.00 KiB	11.33 KiB	67011933	64.79 GiB	1.01 KiB	64.79 GiB	100.00
@@datadir/ib_logfile1	0	0 bytes	0 bytes	66779384	64.70 GiB	1.02 KiB	64.70 GiB	100.00

Profiling General Instance Activity

- Analyzing User Activity
- All event_* summaries are give a number of dimensions to view data
- To analyze connection activity you can do this in 3 ways
 - By User
 - By Host
 - By Account (user@host)
- The following examples are by user, but could be replaced with the host, or account, summary views

Profiling General Instance Activity

- Top Users by IO Latency (or any other event class really) -
events_waits_summary_by_user_by_event_name

```
mysql> SELECT user, SUM(count_star) AS count,  
->      format_time(SUM(sum_timer_wait)) as total_latency  
-> FROM performance_schema.events_waits_summary_by_user_by_event_name  
-> WHERE event_name LIKE 'wait/io/file/%'  
->      AND user IS NOT NULL  
-> GROUP BY user  
-> ORDER BY SUM(sum_timer_wait) DESC;
```

user	count	total_latency
mem3	3046722930	00:41:57.51
service_manager	7131542	00:09:08.22
mark	234	1.43 ms

Replace with other
event classes here

Profiling General Instance Activity

- Top Users by Statement Latency -
events_statements_summary_by_user_by_event_name

```
mysql> SELECT user,  
->      SUM(count_star) AS total_statements,  
->      format_time(SUM(sum_timer_wait)) AS total_latency,  
->      format_time(SUM(sum_timer_wait) / SUM(count_star)) AS avg_latency  
-> FROM performance_schema.events_statements_summary_by_user_by_event_name  
-> WHERE user IS NOT NULL  
-> GROUP BY user  
-> ORDER BY SUM(sum_timer_wait) DESC  
-> LIMIT 5;
```

user	total_statements	total_latency	avg_latency
mem3	1167229105	193.17h	595.78 us
root	2104	00:06:11.61	176.62 ms
service_manager	313171	00:05:05.65	975.97 us
agent_limit	495270	00:03:51.77	467.96 us
mark	227	801.68 ms	3.53 ms

Profiling General Instance Activity

- Top Users by Connections – users, hosts, accounts tables

```
mysql> SELECT * FROM users WHERE user IS NOT NULL ORDER BY current_connections DESC;
```

USER	CURRENT_CONNECTIONS	TOTAL_CONNECTIONS
mem3	24	5956
service_manager	7	10772
mark	1	1

```
3 rows in set (0.07 sec)
```

```
mysql> SELECT * FROM accounts WHERE user IS NOT NULL ORDER BY current_connections DESC;
```

USER	HOST	CURRENT_CONNECTIONS	TOTAL_CONNECTIONS
mem3	[REDACTED].oracle.com	24	5956
service_manager	[REDACTED].oracle.com	6	10771
service_manager	localhost	1	1
mark	[REDACTED].vpn.oracle.com	1	1

Profiling General Instance Activity

- Analyzing Table Activity
 - 3 summary views have been added in 5.6, to compliment Table IO
 - objects_summary_global_by_type
 - A high level aggregate view of database object latency
 - table_io_waits_summary_by_table
 - A detailed aggregate by table, breaking down reads, writes, etc.
 - table_io_waits_summary_by_index_usage
 - Further breaking down usage per index

Profiling General Instance Activity

- Top Tables by Latency - objects_summary_global_by_type

```
mysql> SELECT object_schema AS db_name,  
->      object_name AS table_name,  
->      count_star AS total_events,  
->      format_time(sum_timer_wait) AS total_latency,  
->      format_time(avg_timer_wait) AS avg_latency,  
->      format_time(max_timer_wait) AS max_latency  
-> FROM performance_schema.objects_summary_global_by_type  
-> ORDER BY sum_timer_wait DESC LIMIT 5;
```

Less events, yet higher latency
is a sign that there could be
concurrency issues

db_name	table_name	total_events	total_latency	avg_latency	max_latency
mem23tyr55	inventory_instance_attributes	65321697	00:22:36.90	20.77 us	879.49 ms
mem23tyr55	dc_p_string	86563785	00:13:50.52	9.59 us	324.01 ms
mem23tyr55	dc_p_long	18174843	00:09:44.27	32.15 us	35.07 ms
mem__inventory	MysqlServer	734370	23.70 s	32.27 us	111.50 ms
mem23tyr55	dc_p_double	717954	21.09 s	29.37 us	34.84 ms

Profiling General Instance Activity

- Table Breakdown by Usage - table_io_waits_summary_by_table

```
mysql> SELECT object_schema, object_name,  
->      count_fetch as selects, format_time(sum_timer_fetch) as select_latency,  
->      count_insert as inserts, format_time(sum_timer_insert) as insert_latency,  
->      count_update as updates, format_time(sum_timer_update) as update_latency,  
->      count_delete as deletes, format_time(sum_timer_delete) as delete_latency  
-> FROM performance_schema.table_io_waits_summary_by_table  
-> ORDER BY sum_timer_wait DESC LIMIT 5;
```

object_schema	object_name	selects	select_latency	inserts	insert_latency	updates	update_latency	deletes	delete_latency
mem__inventory	MysqlServer	1251460	43.41 s	51	22.44 ms	167136	15.02 s	0	0 ps
mem__events	events	293364	5.38 s	16128	2.79 s	277428	20.23 s	0	0 ps
mem__events	subjects	1968949	16.38 s	322	2.17 s	0	0 ps	0	0 ps
mem__inventory	Os	1093048	12.70 s	30	4.51 ms	42419	2.84 s	0	0 ps
mem__inventory	Network_v4Addresses	137505	916.01 ms	63768	9.42 s	0	0 ps	63665	1.97 s

Profiling General Instance Activity

- **Tables with Full Scans**
- Search in the “table_io_waits_summary_by_index_usage” table **where index_name is null**
- This is a catch all row for rows read without indexes

```
mysql> SELECT object_schema,  
->      object_name,  
->      count_read AS rows_full_scanned  
-> FROM performance_schema.table_io_waits_summary_by_index_usage  
-> WHERE index_name IS NULL  
->      AND count_read > 0  
-> ORDER BY count_read DESC LIMIT 5;
```

object_schema	object_name	rows_full_scanned
mem23tyr55	rule_alarms	5207641
mem23tyr55	dc_p_string	3989361
mem__inventory	MysqlServer	153463
mem__advisors	advisor_schedules	99974
mem__inventory	Agent	43464

Profiling General Instance Activity

■ Tables with Unused Indexes

- Search for indexes in the “table_io_waits_summary_by_index_usage” **where count_star = 0**
- *Should ensure that you have a representative time frame before using this!*

```
mysql> SELECT object_schema,  
->      object_name,  
->      index_name  
-> FROM performance_schema.table_io_waits_summary_by_index_usage  
-> WHERE index_name IS NOT NULL  
->      AND count_star = 0  
-> ORDER BY object_schema, object_name limit 5;
```

object_schema	object_name	index_name
mem	dc_p_double	PRIMARY
mem	dc_p_double	end_time
mem	dc_p_long	PRIMARY
mem	dc_p_long	end_time
mem	dc_p_string	begin_time

Profiling General Instance Activity

■ host_cache (New in 5.6)

- Exposes hosts cached for DNS
- Counts errors, per error type
- Shows when errors started happening
- Compare *sum_connect_errors* to *max_connect_errors*

```
mysql> DESC host_cache;
```

Field	Type
IP	varchar(64)
HOST	varchar(255)
HOST_VALIDATED	enum('YES','NO')
SUM_CONNECT_ERRORS	bigint(20)
COUNT_HOST_BLOCKED_ERRORS	bigint(20)
COUNT_NAMEINFO_TRANSIENT_ERRORS	bigint(20)
COUNT_NAMEINFO_PERMANENT_ERRORS	bigint(20)
COUNT_FORMAT_ERRORS	bigint(20)
COUNT_ADDRINFO_TRANSIENT_ERRORS	bigint(20)
COUNT_ADDRINFO_PERMANENT_ERRORS	bigint(20)
COUNT_FCRDNS_ERRORS	bigint(20)
COUNT_HOST_ACL_ERRORS	bigint(20)
COUNT_NO_AUTH_PLUGIN_ERRORS	bigint(20)
COUNT_AUTH_PLUGIN_ERRORS	bigint(20)
COUNT_HANDSHAKE_ERRORS	bigint(20)
COUNT_PROXY_USER_ERRORS	bigint(20)
COUNT_PROXY_USER_ACL_ERRORS	bigint(20)
COUNT_AUTHENTICATION_ERRORS	bigint(20)
COUNT_SSL_ERRORS	bigint(20)
COUNT_MAX_USER_CONNECTIONS_ERRORS	bigint(20)
COUNT_MAX_USER_CONNECTIONS_PER_HOUR_ERRORS	bigint(20)
COUNT_DEFAULT_DATABASE_ERRORS	bigint(20)
COUNT_INIT_CONNECT_ERRORS	bigint(20)
COUNT_LOCAL_ERRORS	bigint(20)
COUNT_UNKNOWN_ERRORS	bigint(20)
FIRST_SEEN	timestamp
LAST_SEEN	timestamp
FIRST_ERROR_SEEN	timestamp
LAST_ERROR_SEEN	timestamp

```
29 rows in set (0.02 sec)
```

Profiling General Instance Activity

- Computing Rates of Change
- Looking at the summary data gives a great idea of overall loads
- Looking at the rates of change for certain events can give you a better idea of **utilization** in certain cases
- For example, the ***wait/synch/cond/sql/MYSQL_RELAY_LOG::update_cond*** event tracks the time spent waiting in the replication SQL thread for new events to be written to the relay log – inversely, this is the SQL thread idle time

Profiling General Instance Activity

Slave Load Average


- An example of monitoring this over time
- Given the history, we can also compute *moving averages*

```
mysql> select tstamp, busy_pct, one_min_avg, five_min_avg, fifteen_min_avg
-> from slave_sql_load_average
-> order by tstamp desc limit 10;
```

tstamp	busy_pct	one_min_avg	five_min_avg	fifteen_min_avg
2012-07-24 14:00:29	79.94	67.10	66.92	66.25
2012-07-24 14:00:24	39.97	67.10	66.93	66.15
2012-07-24 14:00:19	79.92	67.11	66.92	66.34
2012-07-24 14:00:14	84.31	67.09	66.93	66.24
2012-07-24 14:00:09	38.62	66.98	66.93	66.11
2012-07-24 14:00:04	83.26	67.03	66.95	66.31
2012-07-24 13:59:59	79.30	66.73	66.90	66.19
2012-07-24 13:59:54	39.97	66.81	66.91	66.09
2012-07-24 13:59:49	79.94	66.81	66.91	66.28
2012-07-24 13:59:44	79.95	66.80	66.91	66.18

<http://www.markleith.co.uk/2012/07/24/a-mysql-replication-load-average-with-performance-schema/>

Profiling Statement Activity



Profiling General Instance Activity

- Computing Rates of Change
- To look at how things change over time, record the events you are interested in within a **persistent** history table
- Get current stats from performance schema for a specific thread, the last row of history, compute the ***time delta*** and ***event count*** delta
- If interested in calculating percentages et al, copy the approach shown in my replication load average – you have to take in to account waiting 100% on other events, or just this event, etc.

Profiling General Instance Activity

- Computing Rates of Change Caveats
- Note, this method can not be used generally, you have to be able to correlate it to specific thread(s)
- As latency is per thread, and you have to compare to a known time delta, summary views do not give this to you (you don't get thread counts per interval)
- However, when tied to a specific threads such as the SQL thread, you can compute other events in the same way – file IO latency for example

Profiling Statement Activity



Profiling Statement Activity

- Statement Profiling Options
- The new instrumentation in 5.6 gives many options
- See statements running currently with `events_statements_current`
- Summary views by user, host, account
- Statement history with `events_statements_history*`
 - And link to stages, and waits, with “***nested events***”
- ***A normalized view with events_statements_summary_by_digest***

Profiling Statement Activity

■ `events_statements_current`

- Exposes the currently executing, or last executed (if no current) statement per thread – executing = “timer_end IS NULL”
- Stats are incremented “live” for running statements

```
mysql> desc events_statements_current;
```

Field	Type
THREAD_ID	int(11)
EVENT_ID	bigint(20) unsigned
END_EVENT_ID	bigint(20) unsigned
EVENT_NAME	varchar(128)
SOURCE	varchar(64)
TIMER_START	bigint(20) unsigned
TIMER_END	bigint(20) unsigned
TIMER_WAIT	bigint(20) unsigned
LOCK_TIME	bigint(20) unsigned
SQL_TEXT	longtext
DIGEST	varchar(32)
DIGEST_TEXT	longtext
CURRENT_SCHEMA	varchar(64)
OBJECT_TYPE	varchar(64)
OBJECT_SCHEMA	varchar(64)
OBJECT_NAME	varchar(64)
OBJECT_INSTANCE_BEGIN	bigint(20) unsigned
MYSQL_ERRNO	int(11)
RETURNED_SQLSTATE	varchar(5)
MESSAGE_TEXT	varchar(128)
ERRORS	bigint(20) unsigned
WARNINGS	bigint(20) unsigned
ROWS_AFFECTED	bigint(20) unsigned
ROWS_SENT	bigint(20) unsigned
ROWS_EXAMINED	bigint(20) unsigned
CREATED_TMP_DISK_TABLES	bigint(20) unsigned
CREATED_TMP_TABLES	bigint(20) unsigned
SELECT_FULL_JOIN	bigint(20) unsigned
SELECT_FULL_RANGE_JOIN	bigint(20) unsigned
SELECT_RANGE	bigint(20) unsigned
SELECT_RANGE_CHECK	bigint(20) unsigned
SELECT_SCAN	bigint(20) unsigned
SORT_MERGE_PASSES	bigint(20) unsigned
SORT_RANGE	bigint(20) unsigned
SORT_ROWS	bigint(20) unsigned
SORT_SCAN	bigint(20) unsigned
NO_INDEX_USED	bigint(20) unsigned
NO_GOOD_INDEX_USED	bigint(20) unsigned
NESTING_EVENT_ID	bigint(20) unsigned
NESTING_EVENT_TYPE	enum('STATEMENT', 'STAGE', 'WAIT')

Profiling Statement Activity

- Generating a Detailed Processlist by JOINing to threads

```
mysql> SELECT pps.thread_id AS thd_id,  
-> pps.processlist_id AS conn_id,  
-> IF(pps.name = 'thread/sql/one_connection',  
-> CONCAT(pps.processlist_user, '@', pps.processlist_host),  
-> REPLACE(pps.name, 'thread/', '')) user,  
-> pps.processlist_db AS db,  
-> pps.processlist_command AS command,  
-> pps.processlist_state AS state,  
-> pps.processlist_time AS time,  
-> format_statement(pps.processlist_info) AS current_statement,  
-> IF(esc.timer_wait IS NOT NULL,  
-> format_statement(esc.sql_text),  
-> NULL) AS last_statement,  
-> IF(esc.timer_wait IS NOT NULL,  
-> format_time(esc.timer_wait),  
-> NULL) AS last_statement_latency,  
-> format_time(esc.lock_time) AS lock_latency,  
-> esc.rows_examined,  
-> esc.rows_sent,  
-> esc.rows_affected,  
-> esc.created_tmp_tables AS tmp_tables,  
-> esc.created_tmp_disk_tables AS tmp_disk_tables,  
-> IF(esc.no_good_index_used > 0 OR esc.no_index_used > 0,  
-> 'YES', 'NO') AS full_scan,  
-> ewc.event_name AS last_wait,  
-> IF(ewc.timer_wait IS NULL AND ewc.event_name IS NOT NULL,  
-> 'Still Waiting',  
-> format_time(ewc.timer_wait)) last_wait_latency,  
-> ewc.source  
-> FROM performance_schema.threads AS pps  
-> LEFT JOIN performance_schema.events_waits_current AS ewc USING (thread_id)  
-> LEFT JOIN performance_schema.events_statements_current AS esc USING (thread_id)  
-> ORDER BY pps.processlist_time DESC, last_wait_latency DESC;
```

```
***** 24. row *****  
thd_id: 71454  
conn_id: 71435  
user: mem3@[REDACTED].oracle.com  
db: mysql  
command: Query  
state: executing  
time: 2  
current_statement: SELECT page_type, IF(t ...  
last_statement: NULL  
last_statement_latency: NULL  
lock_latency: 274.00 us  
rows_examined: 0  
rows_sent: 0  
rows_affected: 0  
tmp_tables: 3  
tmp_disk_tables: 1  
full_scan: YES  
last_wait: wait/io/file/myisam/dfile  
last_wait_latency: 531.34 us  
source: mf_iocache.c:1783
```

Writing temp table
to disk

Profiling Statement Activity

- History of completed statements (*last 10 per session by default!*)

```
mysql> select * from events_statements_history limit 1\G
***** 1. row *****
  THREAD_ID: 378012
   EVENT_ID: 249
 END_EVENT_ID: 288
  EVENT_NAME: statement/sql/select
    SOURCE: mysqld.cc:907
  TIMER_START: 1380151755620846131
  TIMER_END: 1380151756820301872
  TIMER_WAIT: 1199455741
   LOCK_TIME: 0
   SQL_TEXT: select * from events_stages_history
    DIGEST: f13b9e15a70df2108f6aa343860734fa
  DIGEST_TEXT: SELECT * FROM events_stages_history
  CURRENT_SCHEMA: performance_schema
   OBJECT_TYPE: NULL
  OBJECT_SCHEMA: NULL
   OBJECT_NAME: NULL
OBJECT_INSTANCE_BEGIN: NULL
   MYSQL_ERRNO: 0
  RETURNED_SQLSTATE: NULL
  MESSAGE_TEXT: NULL
```

```
      ERRORS: 0
    WARNINGS: 0
  ROWS_AFFECTED: 0
    ROWS_SENT: 125
  ROWS_EXAMINED: 125
  CREATED_TMP_DISK_TABLES: 0
  CREATED_TMP_TABLES: 0
    SELECT_FULL_JOIN: 0
  SELECT_FULL_RANGE_JOIN: 0
    SELECT_RANGE: 0
  SELECT_RANGE_CHECK: 0
    SELECT_SCAN: 1
  SORT_MERGE_PASSES: 0
    SORT_RANGE: 0
    SORT_ROWS: 0
    SORT_SCAN: 0
    NO_INDEX_USED: 1
  NO_GOOD_INDEX_USED: 0
  NESTING_EVENT_ID: NULL
  NESTING_EVENT_TYPE: NULL
```

Profiling Statement Activity

▪ Statement Summaries Per User

```
mysql> SELECT *
-> FROM performance_schema.events_statements_summary_by_user_by_event_name
-> WHERE user = 'mem3'
-> AND sum_timer_wait > 0
-> ORDER BY sum_timer_wait DESC LIMIT 1\G
***** 1. row *****
      USER: mem3
      EVENT_NAME: statement/sql/select
      COUNT_STAR: 46110097
      SUM_TIMER_WAIT: 56305193025555000
      MIN_TIMER_WAIT: 17858000
      AVG_TIMER_WAIT: 1221103000
      MAX_TIMER_WAIT: 47860948487000
      SUM_LOCK_TIME: 4326090581000000
      SUM_ERRORS: 26122
      SUM_WARNINGS: 3194
      SUM_ROWS_AFFECTED: 32
      SUM_ROWS_SENT: 59436762
      SUM_ROWS_EXAMINED: 1953654145
SUM_CREATED_TMP_DISK_TABLES: 14262
SUM_CREATED_TMP_TABLES: 60178
SUM_SELECT_FULL_JOIN: 5926
SUM_SELECT_FULL_RANGE_JOIN: 0
SUM_SELECT_RANGE: 17880
SUM_SELECT_RANGE_CHECK: 0
SUM_SELECT_SCAN: 204539
SUM_SORT_MERGE_PASSES: 0
SUM_SORT_RANGE: 0
SUM_SORT_ROWS: 150801
SUM_SORT_SCAN: 3073
SUM_NO_INDEX_USED: 190598
SUM_NO_GOOD_INDEX_USED: 0
```

Summarized
by statement type,
i.e:
statement/sql/select
statement/sql/update

Profiling Statement Activity

▪ Stage Summaries Per User

```
mysql> SELECT user, event_name AS stage,  
->      count_star AS count,  
->      format_time(sum_timer_wait) as total_latency,  
->      format_time(avg_timer_wait) as avg_latency,  
->      format_time(max_timer_wait) as max_latency  
-> FROM performance_schema.events_stages_summary_by_user_by_event_name  
-> WHERE user = 'mem3'  
->      AND sum_timer_wait > 0  
-> ORDER BY sum_timer_wait DESC;
```

user	stage	count	total_latency	avg_latency	max_latency
mem3	stage/sql/Sending data	742258	00:09:57.41	804.85 us	40.43 s
mem3	stage/sql/init	2499216	00:02:23.17	57.29 us	192.33 ms
mem3	stage/sql/statistics	742397	44.26 s	59.61 us	93.55 ms
mem3	stage/sql/freeing items	1503235	19.32 s	12.85 us	4.13 ms
mem3	stage/sql/updating	205527	19.13 s	93.08 us	163.45 ms
mem3	stage/sql/Opening tables	1006357	13.69 s	13.60 us	14.62 ms
mem3	stage/sql/System lock	997725	11.80 s	11.82 us	2.39 ms
mem3	stage/sql/preparing	742361	6.22 s	8.38 us	2.08 ms
mem3	stage/sql/closing tables	1503135	5.81 s	3.87 us	4.10 ms
mem3	stage/sql/optimizing	744163	4.29 s	5.77 us	4.97 ms
mem3	stage/sql/update	97458	3.98 s	40.83 us	175.50 ms
mem3	stage/sql/checking permissions	1303885	3.16 s	2.43 us	2.53 ms

Profiling Statement Activity

- Tracing Statements
- By taking all available data in the “*_history_long” tables, we can build a complete picture of what sessions and statements are doing
- We can **link** the histories of Statements, Stages and Waits using the “**event_id**” and “**nesting_event_id**” columns, which define hierarchy
- *This is only really effective on development/test systems, production systems with concurrency age history very very quickly, and turning the *_history_long tables on with all instrumentation on is probably not recommend*

Profiling Statement Activity

```
mysql> select event_id, event_name, timer_wait, nesting_event_id  
-> from events_waits_history_long where thread_id = 20;
```

event_id	event_name	timer_wait	nesting_event_id
2547	idle	9264793230	2516
2548	wait/io/socket/sql/client_connection	4974480	2516
2551	wait/synch/mutex/sql/THD::LOCK_thd_data	472500	2550
2552	wait/synch/mutex/sql/THD::LOCK_thd_data	90720	2550
2554	wait/synch/rwlock/sql/LOCK_grant	623700	2553
2556	wait/synch/mutex/sql/MDL_map::mutex	181440	2555
2557	wait/synch/rwlock/sql/MDL_lock::rwlock	128520	2555

```
mysql> select event_id, event_name, timer_wait, nesting_event_id from events_stages_history_long where event_id = 2555;
```


event_id	event_name	timer_wait	nesting_event_id
2555	stage/sql/Opening tables	203212000	2549

```
mysql> select event_id, event_name, timer_wait, substring(sql_text, 1, 40) as sql_text  
-> from events_statements_history_long where event_id = 2549;
```

event_id	event_name	timer_wait	sql_text
2549	statement/sql/select	309041000	select event_id, event_name, timer_wait,

This means we can graph their relationships!

http://www.markleith.co.uk/ps_helper/ps_helper-dump_thread_stack/



Profiling Statement Activity

- Statement Digests
- A new summary view that presents aggregated statistics for ***normalized statements***
- Can be used to drill in to your problem statements instead of using the Slow Query Log, online (*and it has more info*)
- Each statement gets an MD5 “DIGEST”, that is also available to link in to events_statements_current / events_statements_history etc. for raw per query statistics, if wanted

Profiling Statement Activity

```
mysql> desc performance_schema.events_statements_summary_by_digest;
```

Field	Type	Null	Key	Default
DIGEST	varchar(32)	YES		NULL
DIGEST_TEXT	longtext	YES		NULL
COUNT_STAR	bigint(20) unsigned	NO		NULL
SUM_TIMER_WAIT	bigint(20) unsigned	NO		NULL
MIN_TIMER_WAIT	bigint(20) unsigned	NO		NULL
AVG_TIMER_WAIT	bigint(20) unsigned	NO		NULL
MAX_TIMER_WAIT	bigint(20) unsigned	NO		NULL
SUM_LOCK_TIME	bigint(20) unsigned	NO		NULL
SUM_ERRORS	bigint(20) unsigned	NO		NULL
SUM_WARNINGS	bigint(20) unsigned	NO		NULL
SUM_ROWS_AFFECTED	bigint(20) unsigned	NO		NULL
SUM_ROWS_SENT	bigint(20) unsigned	NO		NULL
SUM_ROWS_EXAMINED	bigint(20) unsigned	NO		NULL
SUM_CREATED_TMP_DISK_TABLES	bigint(20) unsigned	NO		NULL
SUM_CREATED_TMP_TABLES	bigint(20) unsigned	NO		NULL
SUM_SELECT_FULL_JOIN	bigint(20) unsigned	NO		NULL
SUM_SELECT_FULL_RANGE_JOIN	bigint(20) unsigned	NO		NULL
SUM_SELECT_RANGE	bigint(20) unsigned	NO		NULL
SUM_SELECT_RANGE_CHECK	bigint(20) unsigned	NO		NULL
SUM_SELECT_SCAN	bigint(20) unsigned	NO		NULL
SUM_SORT_MERGE_PASSES	bigint(20) unsigned	NO		NULL
SUM_SORT_RANGE	bigint(20) unsigned	NO		NULL
SUM_SORT_ROWS	bigint(20) unsigned	NO		NULL
SUM_SORT_SCAN	bigint(20) unsigned	NO		NULL
SUM_NO_INDEX_USED	bigint(20) unsigned	NO		NULL
SUM_NO_GOOD_INDEX_USED	bigint(20) unsigned	NO		NULL
FIRST_SEEN	timestamp	NO		0000-00-00 00:00:00
LAST_SEEN	timestamp	NO		0000-00-00 00:00:00

Profiling Statement Activity

- Statement Normalization
- Normalization folds certain constructs of statements
 - Strip whitespace / comments
 - Replace literals with ?, such as “WHERE foo = 1” becomes “WHERE foo = ?”
 - Folds lists of things, such as “IN (1,2,3)” becomes “IN (...)”
 - Folds multi-row inserts, such as “VALUES (1), (2)” becomes “VALUES (?) /*, ... */”

Profiling Statement Activity

```
mysql> SELECT *
-> FROM performance_schema.events_statements_summary_by_digest
-> ORDER BY sum_timer_wait DESC LIMIT 1\G
***** 1. row *****
      DIGEST: 19b2442e880ef7d9a21db222bae81bdd
      DIGEST_TEXT: SELECT page_type , IF ( TABLE_NAME IS NULL OR INSTR ( TA
, - ? ) ) AS TABLE_NAME , IF ( index_name IS NULL , page_type , index_name ) AS inde
tables WHERE variable_name = ? ) , ? ) , compressed_size ) ) AS total_bytes , SUM ( n
ER_PAGE GROUP BY page_type , TABLE_NAME , index_name ORDER BY NULL
      COUNT_STAR: 1312
      SUM_TIMER_WAIT: 52002971063293000
      MIN_TIMER_WAIT: 31717292272000
      AVG_TIMER_WAIT: 39636410871000
      MAX_TIMER_WAIT: 47860948487000
      SUM_LOCK_TIME: 3721470000000
      SUM_ERRORS: 1
      SUM_WARNINGS: 2622
      SUM_ROWS_AFFECTED: 0
      SUM_ROWS_SENT: 609197
      SUM_ROWS_EXAMINED: 2064202829
      SUM_CREATED_TMP_DISK_TABLES: 2624
      SUM_CREATED_TMP_TABLES: 3936
      SUM_SELECT_FULL_JOIN: 0
      SUM_SELECT_FULL_RANGE_JOIN: 0
      SUM_SELECT_RANGE: 0
      SUM_SELECT_RANGE_CHECK: 0
      SUM_SELECT_SCAN: 2624
      SUM_SORT_MERGE_PASSES: 0
      SUM_SORT_RANGE: 0
      SUM_SORT_ROWS: 0
      SUM_SORT_SCAN: 0
      SUM_NO_INDEX_USED: 1312
      SUM_NO_GOOD_INDEX_USED: 0
      FIRST_SEEN: 2012-09-25 11:50:59
      LAST_SEEN: 2012-09-26 14:00:36
```

Has errors/warnings

More rows scanned
Than returned

Lots of temp tables

Doing full table scans

Profiling Statement Activity

▪ Statements with Temporary Tables

```
mysql> SELECT format_statement(DIGEST_TEXT) AS query,  
->      COUNT_STAR AS exec_count,  
->      SUM_CREATED_TMP_TABLES AS in_memory,  
->      SUM_CREATED_TMP_DISK_TABLES AS on_disk,  
->      ROUND(SUM_CREATED_TMP_TABLES / COUNT_STAR) AS avg_per_query,  
->      ROUND((SUM_CREATED_TMP_DISK_TABLES / SUM_CREATED_TMP_TABLES) * 100) AS to_disk_pct  
-> FROM performance_schema.events_statements_summary_by_digest  
-> WHERE SUM_CREATED_TMP_TABLES > 0  
-> ORDER BY SUM_CREATED_TMP_DISK_TABLES DESC, SUM_CREATED_TMP_TABLES DESC LIMIT 5;
```

query	exec_count	in_memory	on_disk	avg_per_query	to_disk_pct
SELECT * FROM (SELECT digest ... irstSeen` , unix_timestamp ...	1299	9093	3897	7	43
SELECT plugin_name FROM inform ... tus = ? ORDER BY plugin_name	2597	2597	2597	1	100
SELECT page_type , IF (TABLE_ ... E , index_name ORDER BY NULL	1298	3894	2596	3	67
SELECT COUNT (*) AS `num_loc ... ocesslist` WHERE `state` = ?	1567	1567	1567	1	100
SELECT COUNT (*) AS `num_lon ... g_query_time AND `state` = ?	1567	1567	1567	1	100

Profiling Statement Activity

▪ Statements with Full Table Scans

```
mysql> SELECT format_statement(DIGEST_TEXT) AS query,  
-> COUNT_STAR AS exec_count,  
-> SUM_NO_INDEX_USED AS no_index_used_count,  
-> ROUND((SUM_NO_INDEX_USED / COUNT_STAR) * 100) no_index_used_pct,  
-> SUM_ROWS_EXAMINED AS rows_scanned  
-> FROM performance_schema.events_statements_summary_by_digest  
-> WHERE DIGEST_TEXT LIKE 'SELECT%'  
-> AND (SUM_NO_INDEX_USED > 0  
-> OR SUM_NO_GOOD_INDEX_USED > 0)  
-> ORDER BY no_index_used_pct DESC, exec_count DESC LIMIT 5;
```

query	exec_count	no_index_used_count	no_index_used_pct	rows_scanned
SELECT `os0_`.`hid` AS `hid4` ... , `os0_`.`hasVersion` AS ...	93459	93459	100	2947520
SELECT `os0_`.`hid` AS `hid4` ... , `os0_`.`hasVersion` AS ...	39532	39532	100	1369092
SELECT CAST (SUM_NUMBER_OF_BY ... WHERE EVENT_NAME = ? LIMIT ?	6550	6550	100	201740
SELECT `os0_`.`hid` AS `hid4` ... , `os0_`.`hasVersion` AS ...	5956	5956	100	113031
SELECT `os0_`.`hid` AS `hid3` ... sVersion23_39_` , `os0_`.`...	5104	5104	100	122016

Wrapping Up

- This is really just scratching the surface, there are 533 instrument types within 5.6, and many more columns that can be used to filter for different views
- There are many other ways to look at the data that I haven't presented here
- I hope this gives you a helping start in how to start looking at the data though!
- All of these views, and more, are available within ps_helper on my site!